

▼ INF 559: Homework 1

Shaoqian Chen

8831737894

```
from google.colab import drive
drive.mount('/content/gdrive', force_remount=True)
```

```
↳ Mounted at /content/gdrive
```

```
import numpy as np
import matplotlib.pyplot as plt
import time
import random
np.set_printoptions(suppress=True) # Suppress scientific notation while printing numbers
%matplotlib inline
```

▼ 2. Plotting latency and bandwidth

a. Open the test file in unbuffered mode

```
testdata = open("gdrive/My Drive/ubuntu-18.04.3-desktop-amd64.iso",mode="rb",buffering=0)
```

```
testdata
```

```
↳ <_io.FileIO name='gdrive/My Drive/ubuntu-18.04.3-desktop-amd64.iso' mode='rb' close=...
```

b. Sequentially read [1, 4, 16, 64, 256, 1024, 41024, 161024, 641024, 2561024, 1024*1024] size of 4KB. Measure the latency for each iteration in terms of wall-clock time.

```
Q2 = testdata
blocks = [1,4,16,64,256,1024,4*1024,16*1024,64*1024,256*1024,1024*1024]
size = 4
len(blocks)
```

```
↳ 11
```

```
wall_clock_seq = []
for i in range(len(blocks)):
    start = time.time()
```

```

start = time.time()
Q2.readline(blocks[i]*4)
end = time.time()
wall_clock_seq.append(end-start)

```

wall_clock_seq

```

[0.00079345703125,
 1.6927719116210938e-05,
 4.8160552978515625e-05,
 0.00013756752014160156,
 6.723403930664062e-05,
 1.049041748046875e-05,
 0.012387275695800781,
 0.014492034912109375,
 0.00021076202392578125,
 0.0008184909820556641,
 0.0014600753784179688]

```

▼ c. Repeat 2b with random reads instead of sequential.

```

rand_pos = random.sample(range(0,11), 11)

```

rand_pos

```

[9, 6, 10, 0, 5, 3, 1, 2, 4, 7, 8]

```

```

wall_clock_rand = []
for j in rand_pos:
    begin = time.time()
    Q2.readline(blocks[j]*4)
    over = time.time()
    wall_clock_rand.append(over-begin)

```

wall_clock_rand

```

[0.007957696914672852,
 0.0001499652862548828,
 0.0002536773681640625,
 7.867813110351562e-06,
 0.0020613670349121094,
 0.00024175643920898438,
 2.9802322387695312e-05,
 9.274482727050781e-05,
 0.0002567768096923828,
 0.0009310245513916016,
 0.000270843505859375]

```

```
wall_clock_seq
```

```
[0.00079345703125,  
 1.6927719116210938e-05,  
 4.8160552978515625e-05,  
 0.00013756752014160156,  
 6.723403930664062e-05,  
 1.049041748046875e-05,  
 0.012387275695800781,  
 0.014492034912109375,  
 0.00021076202392578125,  
 0.0008184909820556641,  
 0.0014600753784179688]
```

d. Plot the latencies measured in 2b and 2c against the number of blocks read. B

- ▼ should appear on the same plot and the number of blocks should be scaled logar
- Briefly describe your observations from this plot.

```
import numpy as np  
import matplotlib.pyplot as plt  
import pandas as pd
```

```
lat_seq = pd.DataFrame(blocks)  
lat_seq.rename(columns={0:'No. of blocks'})  
lat_seq['Latency(seconds)'] = wall_clock_seq  
Lat_Seq = lat_seq.rename(columns={0:'No.of blocks'})  
Lat_Seq
```

```
[0.00079345703125,  
 1.6927719116210938e-05,  
 4.8160552978515625e-05,  
 0.00013756752014160156,  
 6.723403930664062e-05,  
 1.049041748046875e-05,  
 0.012387275695800781,  
 0.014492034912109375,  
 0.00021076202392578125,  
 0.0008184909820556641,  
 0.0014600753784179688]
```

	No.of blocks	Latency(seconds)
0	1	0.000793
1	4	0.000017
2	16	0.000048
3	64	0.000138
4	256	0.000067
5	1024	0.000010
6	4096	0.012387
7	16384	0.014492
8	65536	0.000211
9	262144	0.000818
10	1048576	0.001460

```
#convert random latency data into dataframe
rand_pos
rand_pos_4 = []
for m in rand_pos:
    rand_pos_4.append(4**m)
rand_pos_4
lat_rand = pd.DataFrame(rand_pos_4)
latency = lat_rand.rename(columns={0:'No. of blocks'})
latency['Latency(seconds)_Random'] = wall_clock_rand
LATENCY = latency.sort_values(by=['No. of blocks'])
LATENCY
```



	No. of blocks	Latency(seconds)_Random
3	1	0.000008
6	4	0.000030
7	16	0.000093
5	64	0.000242
8	256	0.000257
4	1024	0.002061
1	4096	0.000150
9	16384	0.000931
10	65536	0.000271
0	262144	0.007958
2	1048576	0.000254

```
LATENCY['Latency(seconds)_Sequential'] = wall_clock_seq
```

```
LATENCY
```



	No. of blocks	Latency(seconds)_Random	Latency(seconds)_Sequential
3	1	0.000008	0.000793
6	4	0.000030	0.000017
7	16	0.000093	0.000048
5	64	0.000242	0.000138
8	256	0.000257	0.000067
4	1024	0.002061	0.000010
1	4096	0.000150	0.012387
9	16384	0.000931	0.014492
10	65536	0.000271	0.000211
0	262144	0.007958	0.000818
2	1048576	0.000254	0.001460

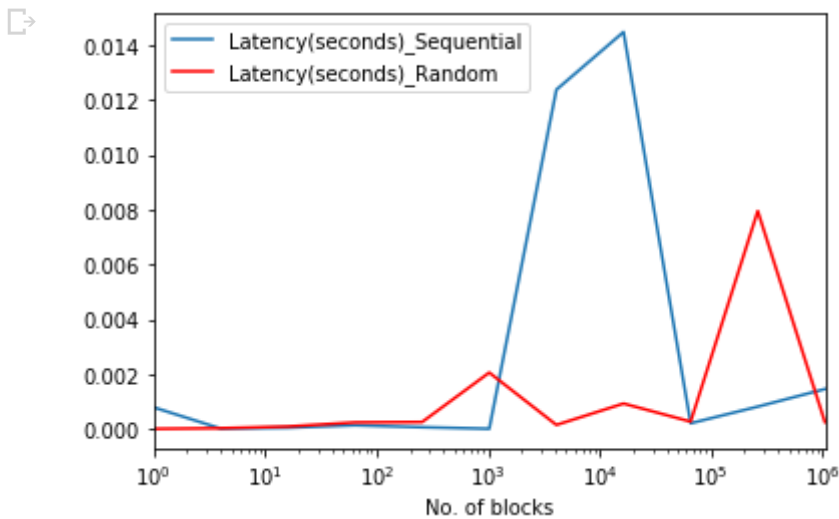
```
test = LATENCY
```

```
test['Blocks'] = ['4^0', '4^1', '4^2', '4^3', '4^4', '4^5', '4^6', '4^7', '4^8', '4^9', '4^10',
test
```



	No. of blocks	Latency(seconds)_Random	Latency(seconds)_Sequential	Blocks
3	1	0.000008	0.000793	4^0
6	4	0.000030	0.000017	4^1
7	16	0.000093	0.000048	4^2
5	64	0.000242	0.000138	4^3
8	256	0.000257	0.000067	4^4
4	1024	0.002061	0.000010	4^5
1	4096	0.000150	0.012387	4^6
9	16384	0.000931	0.014492	4^7
10	65536	0.000271	0.000211	4^8
0	262144	0.007958	0.000818	4^9
2	1048576	0.000254	0.001460	4^10


```
ax = plt.gca()
test.plot(kind='line',x='No. of blocks',y='Latency(seconds)_Sequential',ax=ax)
test.plot(kind='line',x='No. of blocks',y='Latency(seconds)_Random', color='red',ax=ax)
ax.set_xscale('log')
plt.show()
```



e. Calculate the bandwidth for each iteration of 2b and 2c using the latency and the results in the same manner as latency and briefly describe your observations.

```
#Bandwidth is in byte/second
bandwidth = test
```

```
bandwidth_ar = test
bandwidth_df
```




	No. of blocks	Latency(seconds)_Random	Latency(seconds)_Sequential	Blocks
3	1	0.000008	0.000793	4^0
6	4	0.000030	0.000017	4^1
7	16	0.000093	0.000048	4^2
5	64	0.000242	0.000138	4^3
8	256	0.000257	0.000067	4^4
4	1024	0.002061	0.000010	4^5
1	4096	0.000150	0.012387	4^6
9	16384	0.000931	0.014492	4^7
10	65536	0.000271	0.000211	4^8
0	262144	0.007958	0.000818	4^9
2	1048576	0.000254	0.001460	4^10

```
bandwidth_df['Sequential(bandwidth)'] = bandwidth_df['No. of blocks']*4/bandwidth_df[
```

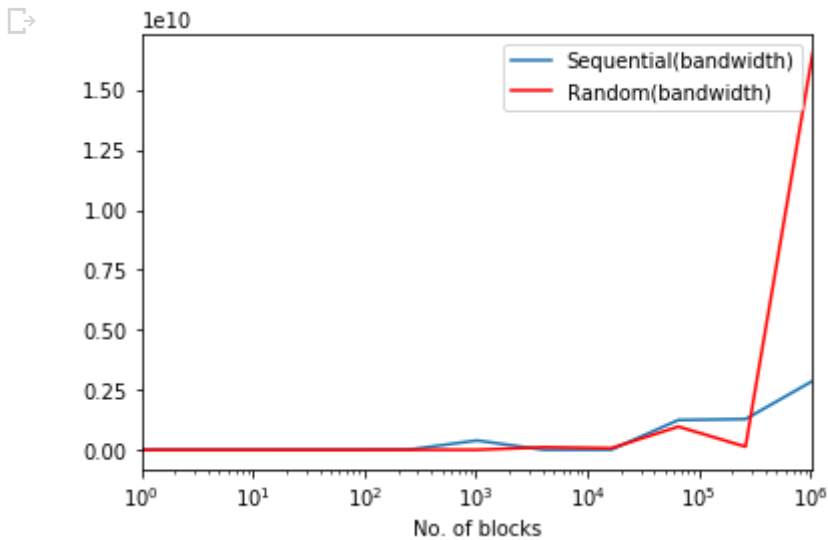
```
bandwidth_df['Random(bandwidth)'] = bandwidth_df['No. of blocks']*4/bandwidth_df['Lat
```

```
bandwidth_df
```



	No. of blocks	Latency(seconds)_Random	Latency(seconds)_Sequential	Blocks
3	1	0.000008	0.000793	4^0
6	4	0.000030	0.000017	4^1
7	16	0.000093	0.000048	4^2
5	64	0.000242	0.000138	4^3
8	256	0.000257	0.000067	4^4
4	1024	0.002061	0.000010	4^5
1	4096	0.000150	0.012387	4^6
9	16384	0.000931	0.014492	4^7
10	65536	0.000271	0.000211	4^8
0	262144	0.007958	0.000818	4^9
2	1048576	0.000254	0.001460	4^10

```
ax = plt.gca()
bandwidth_df.plot(kind='line',x='No. of blocks',y='Sequential(bandwidth)',ax=ax)
bandwidth_df.plot(kind='line',x='No. of blocks',y='Random(bandwidth)', color='red',ax=ax)
ax.set_xscale('log')
plt.show()
```



3. Monte-Carlo simulations

a. Run 10 simulations of steps 2b and 2c and store the results.

```
#2b
"""
For each of the dataframe created below, each column represent a simulation, each
row represents a single iteration within that simulation.

simulation1  simulation2  simulation3...  simulation10
iteration 1
iteration 2
.
.
.
iteration 10
"""

sim_seq = pd.DataFrame()# empty dataframe store all simulation of Sequential case
sim_rand = pd.DataFrame()# empty dataframe store all simulation of Random case
for sim in range(10):
    temp_seq = []
    temp_rand = []
    rand_pos = random.sample(range(0,11), 11)
    for i in range(len(blocks)):
        #add iteration into a temp Sequential list
```



```

start = time.time()
Q2.readline(blocks[i]*4)
end = time.time()
temp_seq.append(end-start)
#add iteration into a temp Random list
begin = time.time()
Q2.readline(blocks[rand_pos[i]]*4)
over = time.time()
temp_rand.append(over-begin)
[temp_rand for _,temp_rand in sorted(zip(rand_pos,temp_rand))]

sim_seq[sim] = temp_seq
sim_rand[sim] = temp_rand
sim_rand

```



	0	1	2	3	4	5	6	7	8
0	0.000009	0.000004	0.000003	0.000523	0.000044	0.000019	0.000088	0.000053	0.000012
1	0.000375	0.000262	0.000071	0.000211	0.000040	0.001903	0.000037	0.000013	0.000077
2	0.000011	0.000155	0.000717	0.000984	0.000005	0.000148	0.000097	0.000002	0.000037
3	0.000136	0.000066	0.000375	0.000150	0.000034	0.005462	0.000037	0.000001	0.000089
4	0.000654	0.000006	0.000367	0.000003	0.000124	0.000037	0.000029	0.000016	0.000005
5	0.000005	0.000680	0.000149	0.000083	0.000003	0.000002	0.000036	0.000010	0.000002
6	0.000005	0.000088	0.000418	0.000190	0.000006	0.000002	0.000003	0.000001	0.000024
7	0.000093	0.000366	0.000242	0.001774	0.000599	0.000001	0.000016	0.000002	0.000059
8	0.000411	0.000136	0.000766	0.000038	0.000038	0.000017	0.000037	0.000002	0.000022
9	0.000088	0.000360	0.000110	0.000010	0.002491	0.000036	0.000010	0.000008	0.000003
10	0.000166	0.000070	0.000014	0.000051	0.005741	0.000013	0.000016	0.000001	0.000004

sim_seq



	0	1	2	3	4	5	6	7	8
0	0.000422	0.000019	0.000017	0.000018	0.000016	0.000020	0.000020	0.000036	0.000012
1	0.000021	0.000012	0.000012	0.000012	0.000010	0.000018	0.000002	0.000002	0.000010
2	0.000044	0.000051	0.000006	0.000005	0.000037	0.000044	0.000001	0.000016	0.000002
3	0.000027	0.000159	0.000165	0.000149	0.000042	0.000145	0.000002	0.000036	0.000001
4	0.000006	0.000255	0.000005	0.000134	0.000044	0.000004	0.000001	0.000013	0.000003
5	0.000183	0.002307	0.000079	0.000136	0.000038	0.000003	0.000001	0.000001	0.000052
6	0.000517	0.000078	0.000133	0.000133	0.000027	0.000030	0.000002	0.000028	0.000020
7	0.000016	0.000212	0.000005	0.000068	0.000033	0.000037	0.000011	0.000029	0.000003
8	0.000264	0.000069	0.000005	0.000607	0.001203	0.000013	0.000001	0.000037	0.000003
9	0.000006	0.001641	0.000299	0.000009	0.000101	0.000002	0.000002	0.000029	0.000029
10	0.000440	0.000005	0.000747	0.000028	0.001549	0.000001	0.000004	0.000036	0.000077

Issue Encountered:

"""

* Each iteration of the random process is using different sequence, need to be specific saperately

* Sequential case works great

"""

#PROBLEM SOLVED

☞ '\n* Each iteration of the random process is using different sequence, need to be

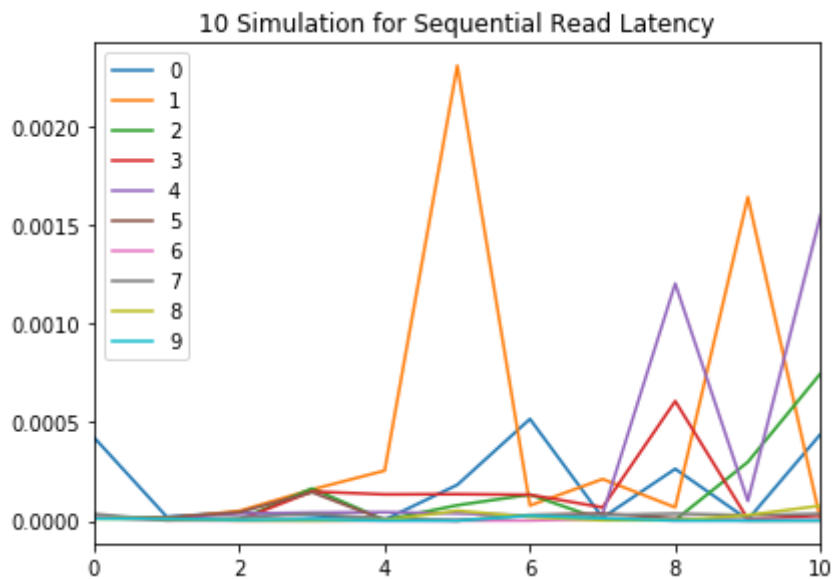
```
sim_seq.plot(subplots=False)
```

```
plt.tight_layout()
```

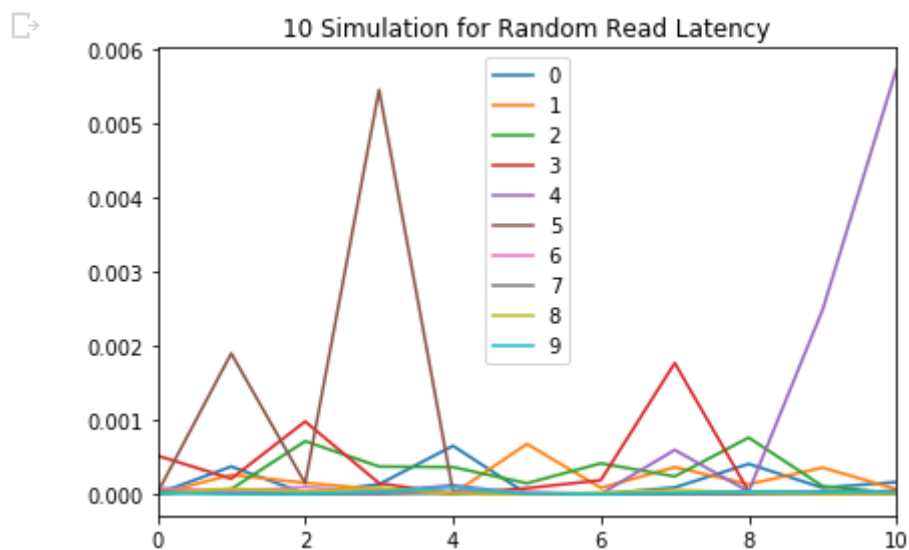
```
plt.title("10 Simulation for Sequential Read Latency")
```

```
plt.show()
```

☞



```
sim_rand.plot(subplots=False)
plt.title("10 Simulation for Random Read Latency")
plt.tight_layout()
plt.show()
```



```
"""
Sort list b according to list a
"""
a = [3,1,2,4,0,12]
b = [0.3,0.1,0.2,0.4,0,1.2]
[b for _,b in sorted(zip(a,b))]

[0, 0.1, 0.2, 0.3, 0.4, 1.2]
```

- ▼ b. For each of the 11 iterations, calculate the mean and standard error4 over the `

```
sim_rand_mean = sim_rand.mean(1)
sim_rand_se = sim_rand.sem(1)
sim_rand_data = pd.DataFrame()
sim_rand_data['Mean'] = sim_rand_mean
sim_rand_data['Standard Error'] = sim_rand_se
```

```
sim_seq_mean = sim_seq.mean(1)
sim_seq_se = sim_seq.sem(1)
sim_seq_data = pd.DataFrame()
sim_seq_data['Mean'] = sim_seq_mean
sim_seq_data['Standard Error'] = sim_seq_se
sim_rand_data
```



	Mean	Standard Error
0	0.000077	0.000050
1	0.000299	0.000182
2	0.000216	0.000109
3	0.000638	0.000537
4	0.000134	0.000068
5	0.000098	0.000066
6	0.000074	0.000043
7	0.000317	0.000173
8	0.000150	0.000078
9	0.000315	0.000244
10	0.000611	0.000570

```
sim_seq_data
```




	Mean	Standard Error
0	0.000059	0.000040
1	0.000011	0.000002
2	0.000021	0.000006
3	0.000073	0.000023
4	0.000047	0.000027
5	0.000280	0.000226
6	0.000099	0.000049
7	0.000043	0.000020
8	0.000220	0.000125
9	0.000212	0.000161
10	0.000289	0.000161

- c. Use the results of 3b to generate errorbar plots for latency and bandwidth. Again, results should be on the same plot and the number of blocks should be scaled logarithmically to your observations from these plots.

```
ax = plt.gca()
plt.errorbar(test['No. of blocks'],sim_seq_data.Mean,yerr=sim_seq_data['Standard Error'])
plt.errorbar(test['No. of blocks'],sim_rand_data.Mean,yerr=sim_rand_data['Standard Error'])
ax.set_xscale('log')
plt.title('Latency Errorbar Plot(10 Simulation)')
plt.legend()
plt.show()
```





```
bandwidth_rand = pd.DataFrame()
for i in range(10):
    bandwidth_rand[i] = test['No. of blocks']*4/sim_rand[i]
```

bandwidth_rand



	0	1	2	3	4	5	6	7	8	9	10
0	1.188661e+11	2.443359e+11	3.998224e+11	2.003666e+09	2.364541e+10	5.567147e+10	1.188661e+11	2.443359e+11	3.998224e+11	2.003666e+09	2.364541e+10
1	4.368689e+07	6.241551e+07	2.298310e+08	7.773696e+07	4.090445e+08	8.610384e+06	4.368689e+07	6.241551e+07	2.298310e+08	7.773696e+07	4.090445e+08
2	3.665039e+11	2.710660e+10	5.846522e+09	4.262706e+09	8.796093e+11	2.837449e+10	3.665039e+11	2.710660e+10	5.846522e+09	4.262706e+09	8.796093e+11
3	2.948544e+04	6.056757e+04	1.066574e+04	2.667284e+04	1.165084e+05	7.323737e+02	2.948544e+04	6.056757e+04	1.066574e+04	2.667284e+04	1.165084e+05
4	6.258604e+06	6.362915e+08	1.114852e+07	1.227134e+09	3.316577e+07	1.094259e+08	6.258604e+06	6.362915e+08	1.114852e+07	1.227134e+09	3.316577e+07
5	5.113056e+07	3.764873e+05	1.715243e+06	3.094357e+06	8.259552e+07	1.342177e+08	5.113056e+07	3.764873e+05	1.715243e+06	3.094357e+06	8.259552e+07
6	3.355443e+06	1.823610e+05	3.826047e+04	8.399107e+04	2.581110e+06	9.586981e+06	3.355443e+06	1.823610e+05	3.826047e+04	8.399107e+04	2.581110e+06
7	6.882960e+05	1.746490e+05	2.644684e+05	3.608488e+04	1.067762e+05	5.368709e+07	6.882960e+05	1.746490e+05	2.644684e+05	3.608488e+04	1.067762e+05
8	2.494174e+06	7.521834e+06	1.336747e+06	2.701237e+07	2.684355e+07	6.049250e+07	2.494174e+06	7.521834e+06	1.336747e+06	2.701237e+07	2.684355e+07
9	7.409108e+08	1.819179e+08	5.936888e+08	6.544712e+09	2.630914e+07	1.832519e+09	7.409108e+08	1.819179e+08	5.936888e+08	6.544712e+09	2.630914e+07
10	1.582031e+09	3.739835e+09	1.928968e+10	5.162026e+09	4.566457e+07	2.036133e+10	1.582031e+09	3.739835e+09	1.928968e+10	5.162026e+09	4.566457e+07

```
bandwidth_seq = pd.DataFrame()
for j in range(10):
    bandwidth_seq[j] = test['No. of blocks']*4/sim_seq[j]
bandwidth_seq
```



	0	1	2	3	4	5	
0	2.481968e+09	5.638521e+10	6.194432e+10	5.864062e+10	6.467715e+10	5.298851e+10	5.2
1	7.898790e+08	1.402438e+09	1.321528e+09	1.347441e+09	1.561806e+09	9.162597e+08	7.6
2	9.509290e+10	8.259242e+10	6.515624e+11	8.377231e+11	1.127704e+11	9.613216e+10	2.9
3	1.458888e+05	2.511559e+04	2.431481e+04	2.688656e+04	9.478653e+04	2.750363e+04	2.3
4	7.158279e+08	1.607097e+07	8.589935e+08	3.046076e+07	9.286416e+07	1.010581e+09	2.8
5	1.399924e+06	1.109581e+05	3.224450e+06	1.887068e+06	6.795834e+06	9.761289e+07	2.7
6	3.096856e+04	2.046002e+05	1.200516e+05	1.206994e+05	5.835553e+05	5.284163e+05	9.5
7	3.890369e+06	3.019521e+05	1.412818e+07	9.353152e+05	1.917396e+06	1.731842e+06	5.9
8	3.883334e+06	1.491308e+07	2.147484e+08	1.686285e+06	8.511628e+05	7.953643e+07	7.7
9	1.018066e+10	3.992997e+07	2.193758e+08	7.429133e+09	6.467715e+08	3.926827e+10	2.7
10	5.962644e+08	4.997780e+10	3.509453e+08	9.397535e+09	1.692859e+08	1.832519e+11	6.4

```
bandwidth_rand_mean = bandwidth_rand.mean(1)
bandwidth_rand_se = bandwidth_rand.sem(1)
bandwidth_rand_data = pd.DataFrame()
bandwidth_rand_data['Mean'] = bandwidth_rand_mean
bandwidth_rand_data['Standard Error'] = bandwidth_rand_se
```

```
bandwidth_seq_mean = bandwidth_seq.mean(1)
bandwidth_seq_se = bandwidth_seq.sem(1)
bandwidth_seq_data = pd.DataFrame()
bandwidth_seq_data['Mean'] = bandwidth_seq_mean
bandwidth_seq_data['Standard Error'] = bandwidth_seq_se
bandwidth_rand_data
```



	Mean	Standard Error
0	1.053671e+11	3.981219e+10
1	7.640423e+08	4.747678e+08
2	6.494658e+11	3.221172e+11
3	3.319265e+05	2.741738e+05
4	3.203019e+08	1.307934e+08
5	4.512527e+07	1.588971e+07
6	4.588615e+06	1.591553e+06
7	1.022802e+07	6.091596e+06
8	8.405579e+07	5.912490e+07
9	4.796002e+09	2.056154e+09
10	3.177794e+10	1.775003e+10

bandwidth_seq_data



	Mean	Standard Error
0	5.571436e+10	8.077423e+09
1	2.577028e+09	8.473206e+08
2	8.413756e+11	3.293869e+11
3	6.716192e+05	3.770713e+05
4	9.593658e+08	3.146505e+08
5	6.630763e+07	2.734945e+07
6	1.314165e+06	9.228664e+05
7	5.842941e+06	2.227603e+06
8	2.030637e+08	8.551019e+07
9	1.447540e+10	6.140918e+09
10	4.564976e+10	2.067349e+10

```
ax = plt.gca()
plt.errorbar(test['No. of blocks'],bandwidth_seq_data.Mean,yerr=bandwidth_seq_data['S
plt.errorbar(test['No. of blocks'],bandwidth_rand_data.Mean,yerr=bandwidth_rand_data[
ax.set_xscale('log')
plt.title('Bandwidth Errorbar Plot(10 Simulation)')
plt.legend()
plt.show()
```