

# Data Models

# Critical aspects of data

- Semantics
  - What do the data mean?
  - Ontology, taxonomy, schema
- Syntax
  - How are the data structured
  - XML, HDF, CSV

# What is a data model?

- Structure of the data
- Operations on the data
- Constraints on the data
  
- Types of data models
  - Structured data
  - Semi-structured

# Data Modeling/Design

- Conceptual design
  - Model independent of the choice of implementation, concepts and relationships
  - Create a model of a “mini-world”
- Logical Design
  - Full details of data without regard to physical implementation
    - all attributes, types, relationships, cardinality, etc...
- Physical Data Model
  - Mapping into specific implementation details (DBMS, XML, etc).

# Steps in Building a DB Application

- Step 0: pick an application domain
  - E.g., course management
- Step 1: conceptual design
  - Decide on what to model in the application domain
    - E.g., instructors, students, courses, etc.
  - need a modeling language to express what you want
  - ER model is the most popular such language
  - output: an ER diagram of the app. domain

# Steps in Building a DB Application

- Step 2: pick a type of DBMS
  - Here we use relational DBMS
- Step 3: translate ER design to a relational schema
  - use a set of rules to translate ER to rel. schema
  - use a set of schema refinement rules to transform the above rel. schema into a **good** rel. schema
- At this point
  - you have a good relational schema on paper

# Steps in Building a DB Application

- Subsequent steps include
  - implement your relational DBMS using a "database programming language" called SQL
  - ordinary users cannot interact with the database directly
  - and the database also cannot do everything you want
  - hence write your application program in Php, C++, Java, Python, etc. to handle the interaction and take care of things that the database cannot do
- So, the first thing we should start with is to learn ER model ...

# ER Model

- Gives us a language to specify
  - what information the db must hold
  - what are the relationships among components of that information
- Proposed by Peter Chen in 1976
- What we will cover
  - basic stuff
  - subclasses
  - constraints
  - weak entity sets
  - design principles




<http://www.csc.lsu.edu/~chen/>



# Entity Relationship Models

- The E/R model allows us to sketch data model designs (schema).
  - Includes some constraints, but not operations.
- Designs are pictures called *entity-relationship diagrams*.

# A Simple ER Model




Mini-world of instructors and courses

In English:

- Instructors teach courses
- Courses have instructors
- Instructors have a name and phone number
- Courses have a name and a room number

## One possible instance...



What about the second section?

# Entity-Relationship Model

- Semantic model
- We model concepts
  - Above abstraction in database
- Identify set of entities and relationships between those entities

# Entity Sets

- *Entity* = “thing” or object.
  - What are the set of things we care about? May be physical (person) or conceptual (e.g. vacation)
  - Must be possible to distinguish one entity from another
- *Entity set* = collection of similar entities.
  - Similar to a class in object-oriented languages.


# Attributes

- A fact, aspect, or property of (the entities of) an entity set.
  - Attributes are simple values, e.g. integers or character strings, not structs, sets, etc.
- Attributes have
  - A name
  - An associated entity
  - Domains of possible values
  - Values from the domain for each instance of the entity

# E/R Diagrams

- In an entity-relationship diagram:
  - Entity set = rectangle.
  - Attribute = oval, with a line to the rectangle representing its entity set.

## Example:



- Entity set Beers has two attributes, name and manf (manufacturer).
- Each Beers entity has values for these two attributes, e.g. (Bud, Anheuser-Busch)



# Entity or Attribute?

- How do you choose?
  - They both represent objects or facts about the world
  - They both often capture “nouns:
- General guidelines
  - Entity can have smaller subparts, attributes cannot
  - Entities can have relationships, but attributes cannot


# Entity or Attribute?

- Beer
- Supplier
- Manufacturer
- Bar
- Drinker
- License
- Address
- Price
- Bottle size

# Relationships

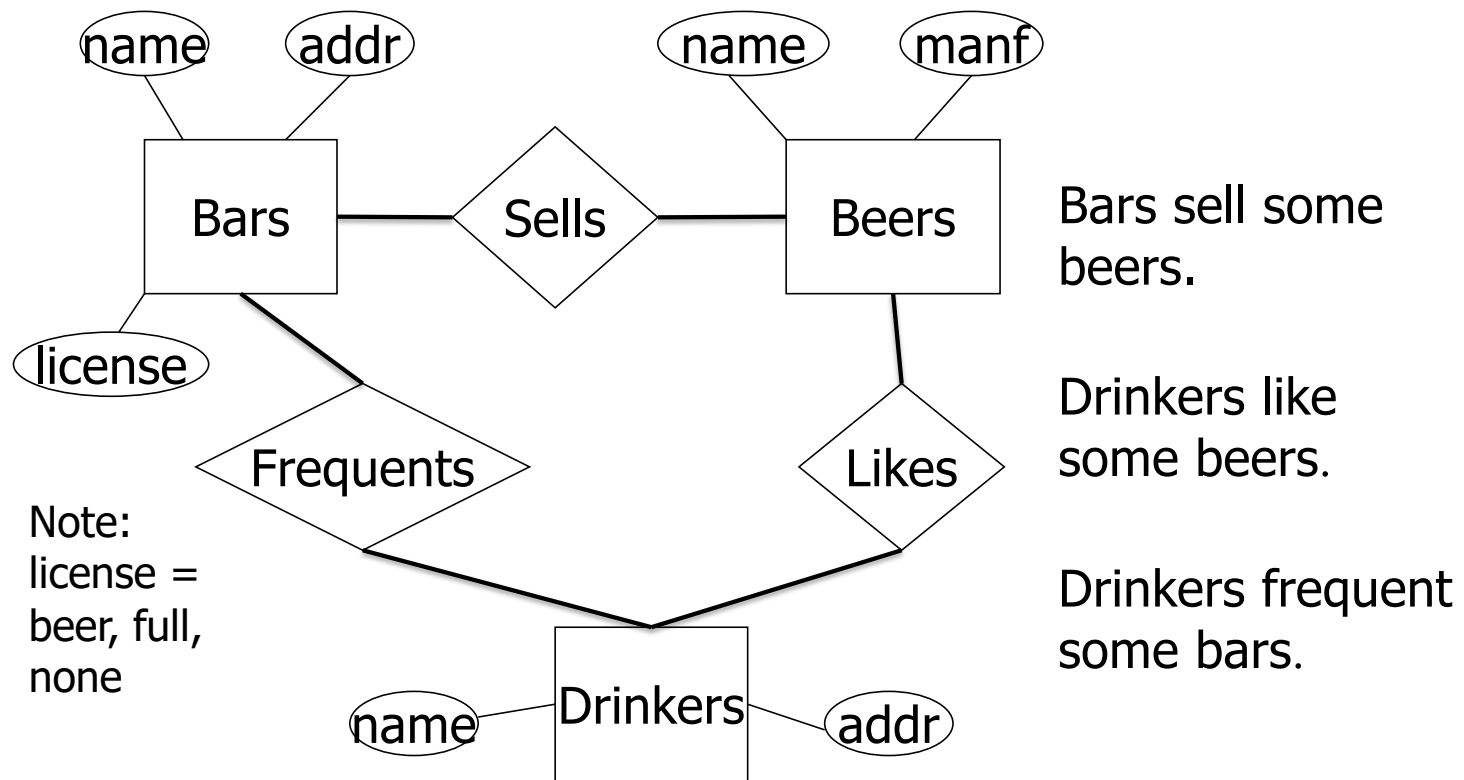
- A **relationship** connects two or more entity sets.
- It is represented by a diamond, with lines to each of the entity sets involved.
- Relations have
  - A name
  - A set of entities that participate in them
  - A degree
  - A cardinality ratio

## Example: Relationships



Note:  
license = beer, full, none

## Example: Relationships



# Relationship Set

- The current “value” of an entity set is the set of entities that belong to it.
  - **Example**: the set of all bars in our database.
- The “value” of a relationship is a *relationship set*, a set of tuples with one component for each related entity set.

## Example: Relationship Set

- For the relationship **Sells**, we might have a relationship set like:

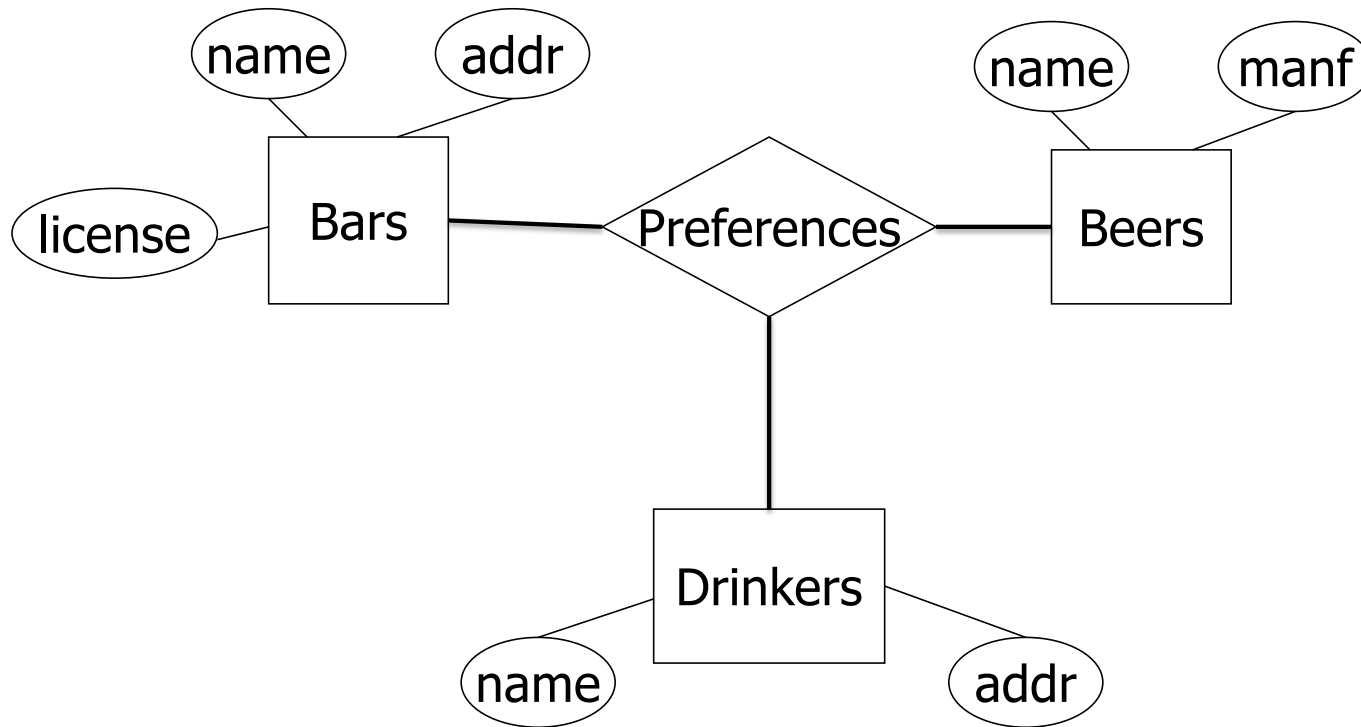
Bar	Beer
Joe' s Bar	Bud
Joe' s Bar	Miller
Sue' s Bar	Bud
Sue' s Bar	Pete' s Ale
Sue' s Bar	Bud Lite

# Multiway Relationships


- Sometimes, we need a relationship that connects more than two entity sets.
- Suppose that drinkers will only drink certain beers at certain bars.
  - Our three binary relationships **Likes**, **Sells**, and **Frequents** do not allow us to make this distinction.
  - But a 3-way relationship would.



## Example: 3-Way Relationship




# Cardinality



# Cardinalities

- No restriction in how often an entity can be in a relationship R
  - Can connect to zero, one, or many
- Mini-world semantics might require more specificity



# Many-Many Relationships


- Focus: **binary** relationships, such as **Sells** between **Bars** and **Beers**.
- In a **many-many relationship**, an entity of either set can be connected to many entities of the other set.
  - E.g., a bar sells many beers; a beer is sold by many bars.



# One-One Relationships

- In a *one-one relationship*, each entity of either entity set is related to at most one entity of the other set.

# In Pictures:



one-one

**Example:** Relationship **Best-seller** between entity sets **Manfs** (manufacturer) and **Beers**.

- A beer cannot be made by more than one manufacturer, and no manufacturer can have more than one best-seller (assume no ties).




31

# Many-One Relationships

- Some binary relationships are *many -one* from one entity set to another.
- Each entity of the first set is connected to at most one entity of the second set.
- But an entity of the second set can be connected to zero, one, or many entities of the first set.
- Cardinality of 1 is on the “many” side!



## In Pictures:



many-one

## Example: Many-One Relationship

- Favorite, from Drinkers to Beers is many-one.
- A drinker has at most one favorite beer.
- But a beer can be the favorite of any number of drinkers, including zero.



34


# One to Many

- Same idea as many-to-one but from opposite perspective:
  - Each entity in the second set is connected to at most one entity in the first

# Representing “Multiplicity”


- Show a many-one relationship by an arrow entering the “one” side.
  - **Remember**: Like a functional dependency.
- Show a one-one relationship by arrows entering both entity sets.
- **Rounded arrow** = “exactly one,” i.e., each entity of the first set is related to exactly one entity of the target set.

## Example: Many-One Relationship








Notice: two relationships connect the same entity sets, but are different.

## In the E/R Diagram



A beer is the best-seller for 0 or 1 manufacturer.

A manufacturer has exactly one best seller.




# Attributes on Relationships

- Sometimes it is useful to attach an attribute to a relationship.
- Think of this attribute as a property of tuples in the relationship set.



## Example: Attribute on Relationship




Price is a function of both the bar and the beer,  
not of one alone.

# Equivalent Diagrams Without Attributes on Relationships

- Create an entity set representing values of the attribute.
- Make that entity set participate in the relationship.

## Example: Removing an Attribute from a Relationship

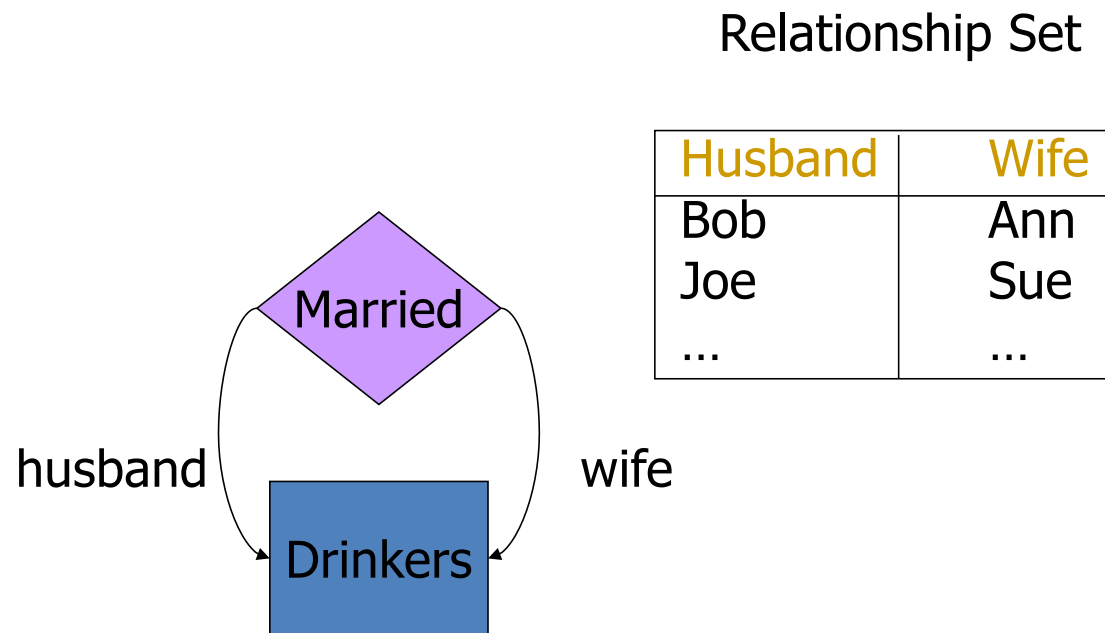


Note convention: arrow from multiway relationship = “all other entity sets together determine a unique one of these.”


# Roles

- Sometimes an entity set appears more than once in a relationship.
- Label the edges between the relationship and the entity set with names called *roles*.

## Example: Roles



## Example: Roles



Relationship Set

Buddy1	Buddy2
Bob	Ann
Joe	Sue
Ann	Bob
Joe	Moe
...	...

# Subclasses


- *Subclass* = special case = fewer entities = more properties.
- *Example*: Ales are a kind of beer.
  - Not every beer is an ale, but some are.
  - Let us suppose that in addition to all the *properties* (attributes and relationships) of beers, ales also have the attribute *color*.

# Subclasses in E/R Diagrams

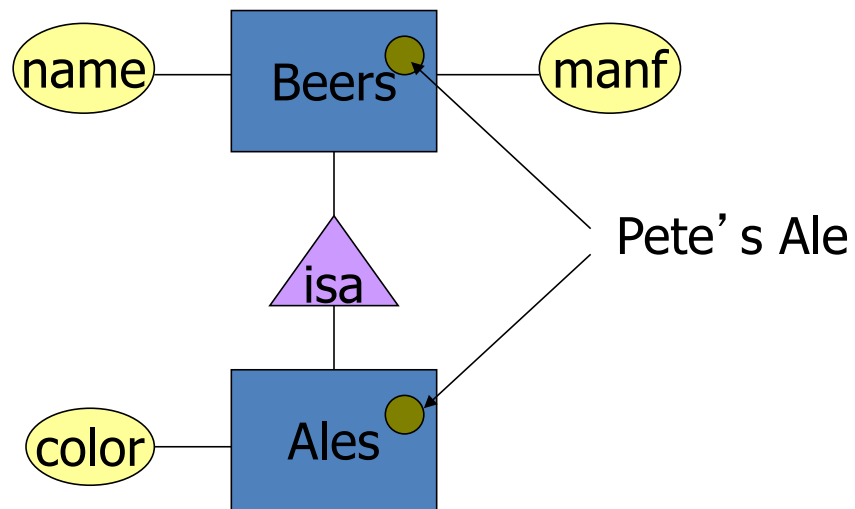
- Assume subclasses form a tree.
  - I.e., no multiple inheritance.
- Isa triangles indicate the subclass relationship.
  - Point to the superclass.



## Example: Subclasses



## Example: Representatives of Entities




# Keys

- A key is a set of attributes for one entity set such that no two entities in this set have the same value for the key.
  - It is allowed for two entities to agree on some, but not all, of the key attributes.
- We must designate a key for every entity set.


## Keys in E/R Diagrams

- Underline the key attribute(s).
- In an Isa hierarchy, only the root entity set has a key, and it must serve as the key for all entities in the hierarchy.

Example: **name** is Key for Beers



## Example: a Multi-attribute Key



- Note that **hours** and **room** could also serve as a key, but we must select only one key.

# Weak Entity Sets


- Occasionally, entities of an entity set need “help” to identify them uniquely.
  - If you have a many to one relationship with another entity, you can use attributes of that entity to form a key
- We call this entity set *weak*

## Example: Weak Entity Set

- **name** is almost a key for football players, but there might be two with the same name.
- **number** is certainly not a key, since players on two teams could have the same number.
- But **number**, together with the team **name** related to the player by **Plays-on** should be unique.



# In E/R Diagrams



Note: must be rounded because each player needs a team to help with the key.

- Double diamond for *supporting* many-one relationship.
- Double rectangle for the weak entity set.

# Weak Entity-Set Rules

- A weak entity set has one or more many-one relationships to other (supporting) entity sets.
  - Not every many-one relationship from a weak entity set need be supporting.
  - But supporting relationships must have a rounded arrow (entity at the “one” end is guaranteed).

## Weak Entity-Set Rules – (2)

- The key for a weak entity set is its own underlined attributes and the keys for the supporting entity sets.
  - E.g., (player) number and (team) name is a key for Players in the previous example.


# Design Techniques

1. Avoid redundancy.
2. Limit the use of weak entity sets.
3. Don't use an entity set when an attribute will do.

# Avoiding Redundancy


- *Redundancy* = saying the same thing in two (or more) different ways.
- Wastes space and (more importantly) encourages inconsistency.
  - Two representations of the same fact become inconsistent if we change one and forget to change the other.
  - Recall anomalies due to FD' s.

## Example: Good




This design gives the address of each manufacturer exactly once.

## Example: Bad



This design states the manufacturer of a beer twice: as an attribute and as a related entity.

## Example: Bad




This design repeats the manufacturer's address once for each beer and loses the address if there are temporarily no beers for a manufacturer.



# Entity Sets Versus Attributes


- An entity set should satisfy at least one of the following conditions:
  - It is more than the name of something; it has at least one nonkey attribute.
  - or
  - It is the “many” in a many-one or many-many relationship.

## Example: Good




- **Manfs** deserves to be an entity set because of the nonkey attribute **addr**.
- **Beers** deserves to be an entity set because it is the “many” of the many-one relationship **ManfBy**.

## Example: Good



There is no need to make the manufacturer an entity set, because we record nothing about manufacturers besides their name.

## Example: Bad



Since the manufacturer is nothing but a name, and is not at the “many” end of any relationship, it should not be an entity set.

## Don' t Overuse Weak Entity Sets

- Beginning database designers often doubt that anything could be a key by itself.
  - They make all entity sets weak, supported by all other entity sets to which they are linked.
- In reality, we usually create unique ID' s for entity sets.
  - Examples include social-security numbers, automobile VIN' s etc.

# When Do We Need Weak Entity Sets?

- The usual reason is that there is no global authority capable of creating unique ID' s.
- **Example:** it is unlikely that there could be an agreement to assign unique player numbers across all football teams in the world.

## Next up....

- Relational databases