

Cloud Storage

Cloud Data Centers....



Early days...

...today



Observations....

- Pervasive failure of system components
 - Especially if you want to use cheap commodity stuff
- Emerging workloads on very large files
 - “Big data,”
 - Multi-gigabyte files, terabyte data sets, much larger than traditional many small file design points
- Random writes are rare
 - Updates are frequently appends
 - More reads than writes, often sequential
- Co-designing application and file system interfaces can have significant benefits
 - E.g. alternative consistency models

Design points

- Use inexpensive commodity parts
 - No expensive disks or controllers
- Small number of large files
 - A few million files, but 100Mb or larger
- Workload consisting of
 - mostly of large streaming reads (100s of KB) and small random reads (few KBs).
 - Large, sequential writes that append
 - Random writes rare
- Many concurrent operations
 - E.g. parallel processing
- Bandwidth is more important than latency

Amazon S3

- Storage on the “cloud”
- Object store with web service interfaces hosted by Amazon


The Simple Storage Service

- **Buckets**
 - Container for objects stored in Amazon S3
- **Objects**
 - Data + Metadata
 - Metadata consists of name/value pairs
- **Keys**
 - A key is the unique identifier for an object within a bucket

Buckets


Every bucket has a name.

Can refer to a bucket by URL:
<http://s3.amazonaws.com/foo>



Objects

Every object has a name.
Must be unique in bucket



Operations on Objects

- Object consists of
 - Key: the object name used to retrieve the object
 - Version ID
 - Value
 - Metadata
- Can store, retrieve and delete objects
 - Retrieve done on whole object
- Objects can be versioned

Representational State Transfer (REST)

- Client-Server
 - a pull-based interaction style: consuming components pull representations.
- Stateless
 - each request from client to server must contain all the information necessary to understand the request, and cannot take advantage of any stored context on the server.
- Cachable
 - to improve network efficiency responses must be capable of being labeled as cacheable or non-cacheable.
- Uniform interface: all resources are accessed with a generic interface
 - (e.g., HTTP GET, POST, PUT, DELETE).
- Layered components
 - intermediaries, such as proxy servers, cache servers, gateways, etc, can be inserted between clients and resources to support performance, security, etc.

CRUD to HTTP

- Create → PUT with a new URI or POST with a base URI returning a new URI
- Read → GET
- Update → PUT,
- Delete → DELETE.

Retrieving and Object

- Object may be retrieved in whole or in parts
- Example:

GET /my-image.jpg HTTP/1.1

Host: bucket.s3.amazonaws.com

Date: Wed, 28 Oct 2009 22:32:00 GMT


Authorization: *authorization string*

Updating an Object

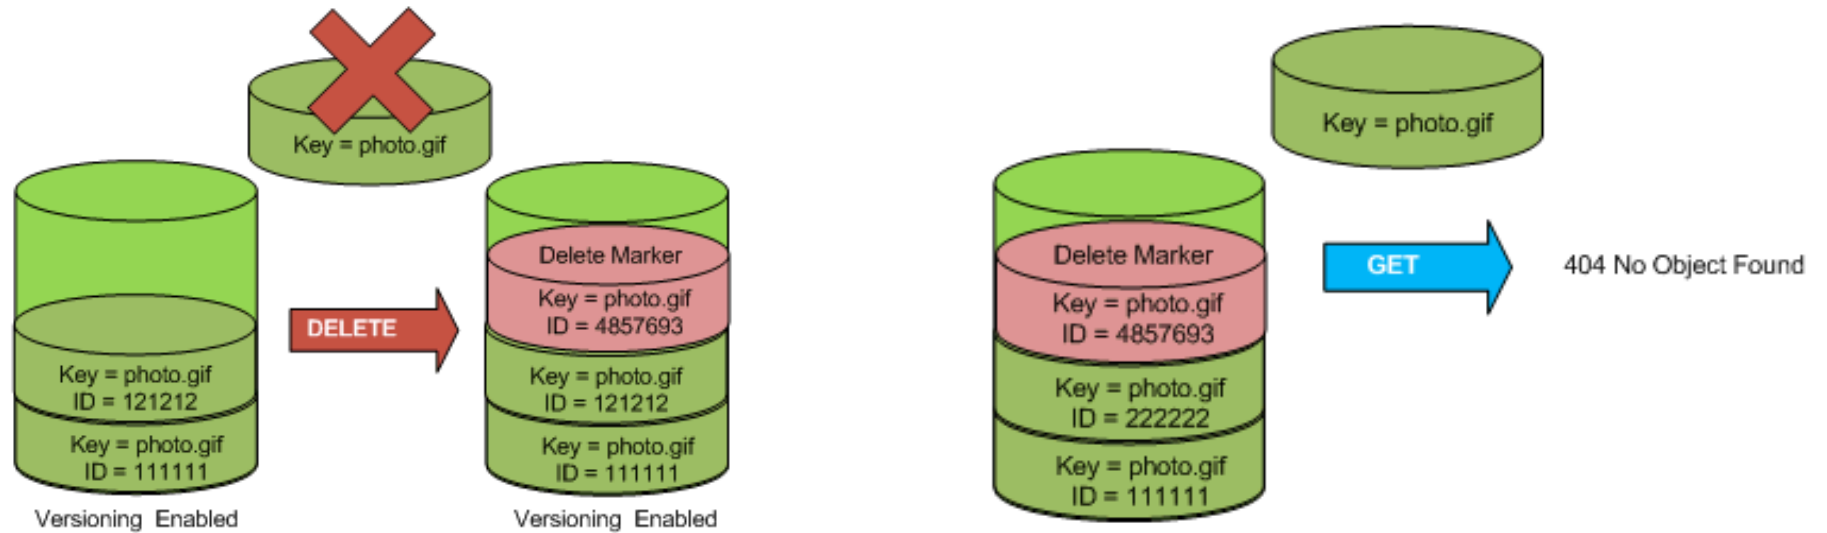
- Only full objects are updated into S3
- Example

```
PUT /my-image.jpg HTTP/1.1
Host: myBucket.s3.amazonaws.com
Date: Wed, 12 Oct 2009 17:50:00 GMT
Authorization: authorization string
Content-Type: text/plain
Content-Length: 11434
Expect: 100-continue
[11434 bytes of object data]
```


Object Versioning



Deleting an Object



Retrieving a Version



Storage Classes

- Standard
- Infrequent Access
- Glacier
- Use cases driven by pricing for storage and access
 - <https://aws.amazon.com/s3/pricing/>

S3 Scalability and Availability

- To achieve reliability, there are many copies of the data kept
 - Just like mirroring in RAID
- However, because of network latency, waiting to update all of the copies is not practical

What is consistency?

- If you have data that may be in multiple locations, and accessed by multiple users, what happens when you update the data?
 - When is a write seen by a read?
- For example:
 - Assume a file is replicated on nodes M and N
 - The client A writes to the to node N
 - After a period of time t , client B reads the file from node M
 - The consistency model has to determine whether client B sees the write from client A or not.

Consistency Models

- *Strong consistency.*
 - After the update completes, any subsequent access (by A, B, or C) will return the updated value.
- *Weak consistency.*
 - The system does not guarantee that subsequent accesses will return the updated value.
- *Eventual consistency.*
 - This is a specific form of weak consistency;
 - the storage system guarantees that if no new updates are made to the object, eventually all accesses will return the last updated value.

Creating New Objects

- Read after write consistency for new objects

PUT /key-prefix/cool-file.jpg 200

GET /key-prefix/cool-file.jpg 200

- But in some cases, not quite...

GET /key-prefix/cool-file.jpg 404


PUT /key-prefix/cool-file.jpg 200

GET /key-prefix/cool-file.jpg 404

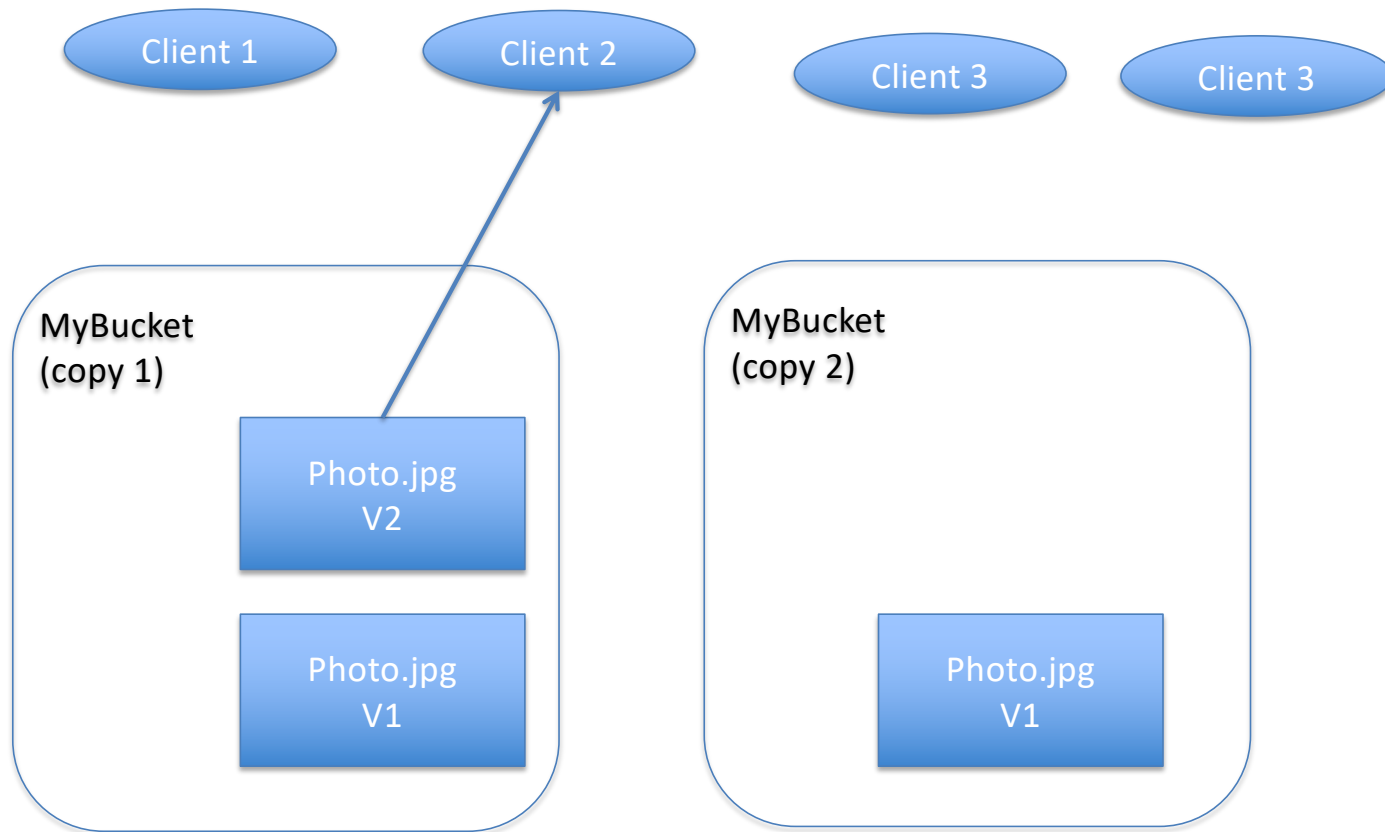
What about updates?

- S3 has “eventual consistency” for updates
- Reads may return old values for an object
- Eventually all reads will return the same value

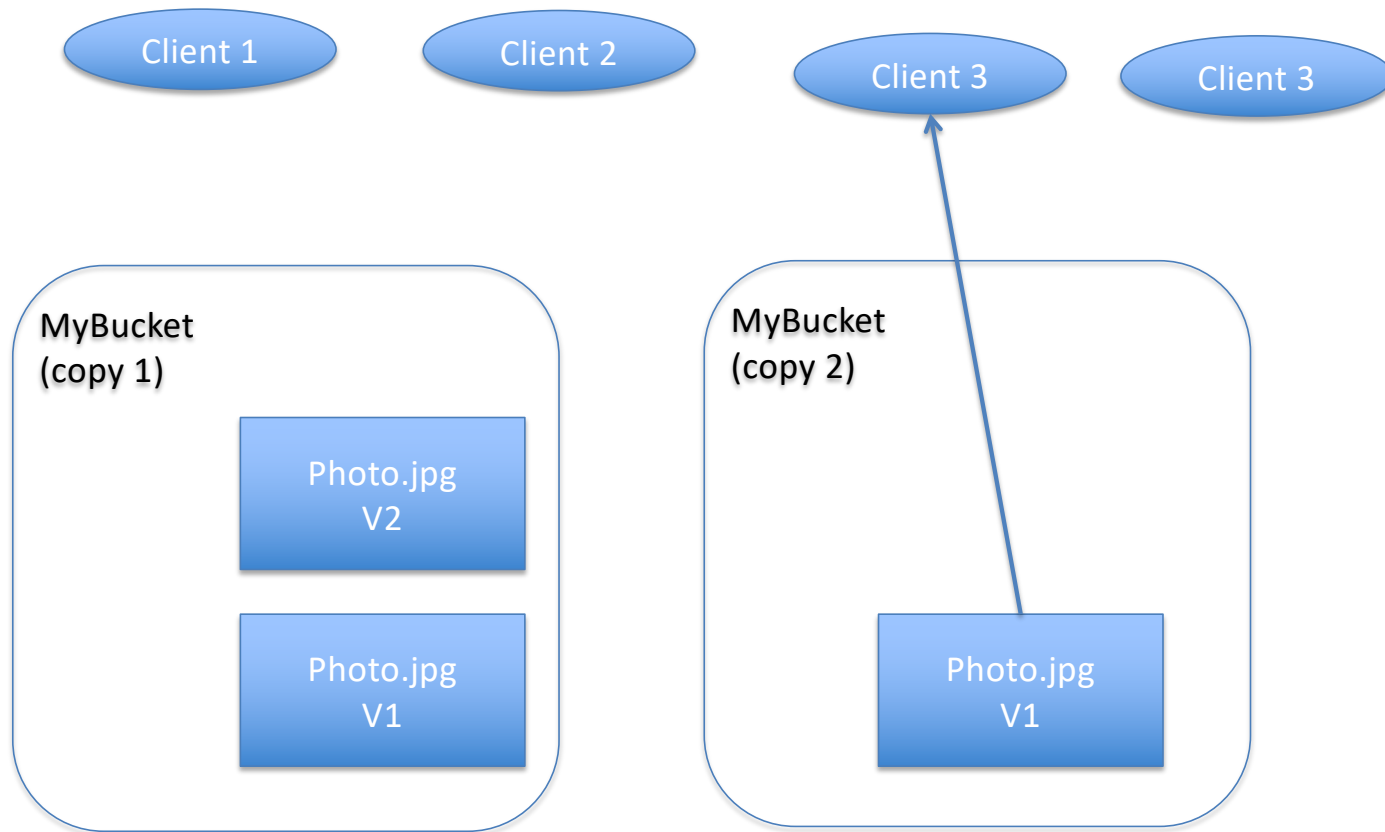
Eventual Consistency



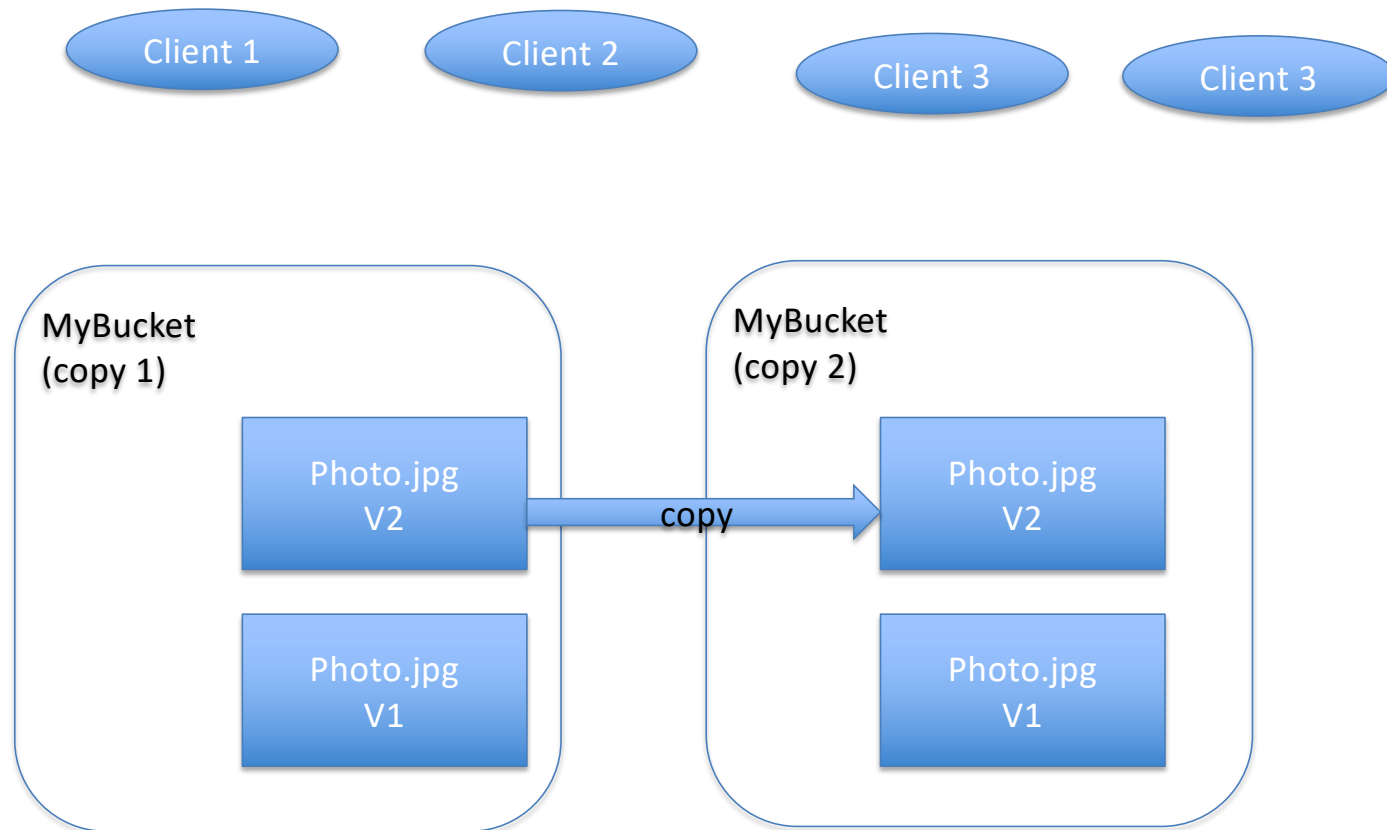
Eventual Consistency



Eventual Consistency



Eventual Consistency



Eventual Consistency

