

Amazon Dynamo DB

- Highly scalable key value store
 - Support for document operations
- Hosted service
 - Focus on ease of configuration and use
- Support for transactions and consistent reads

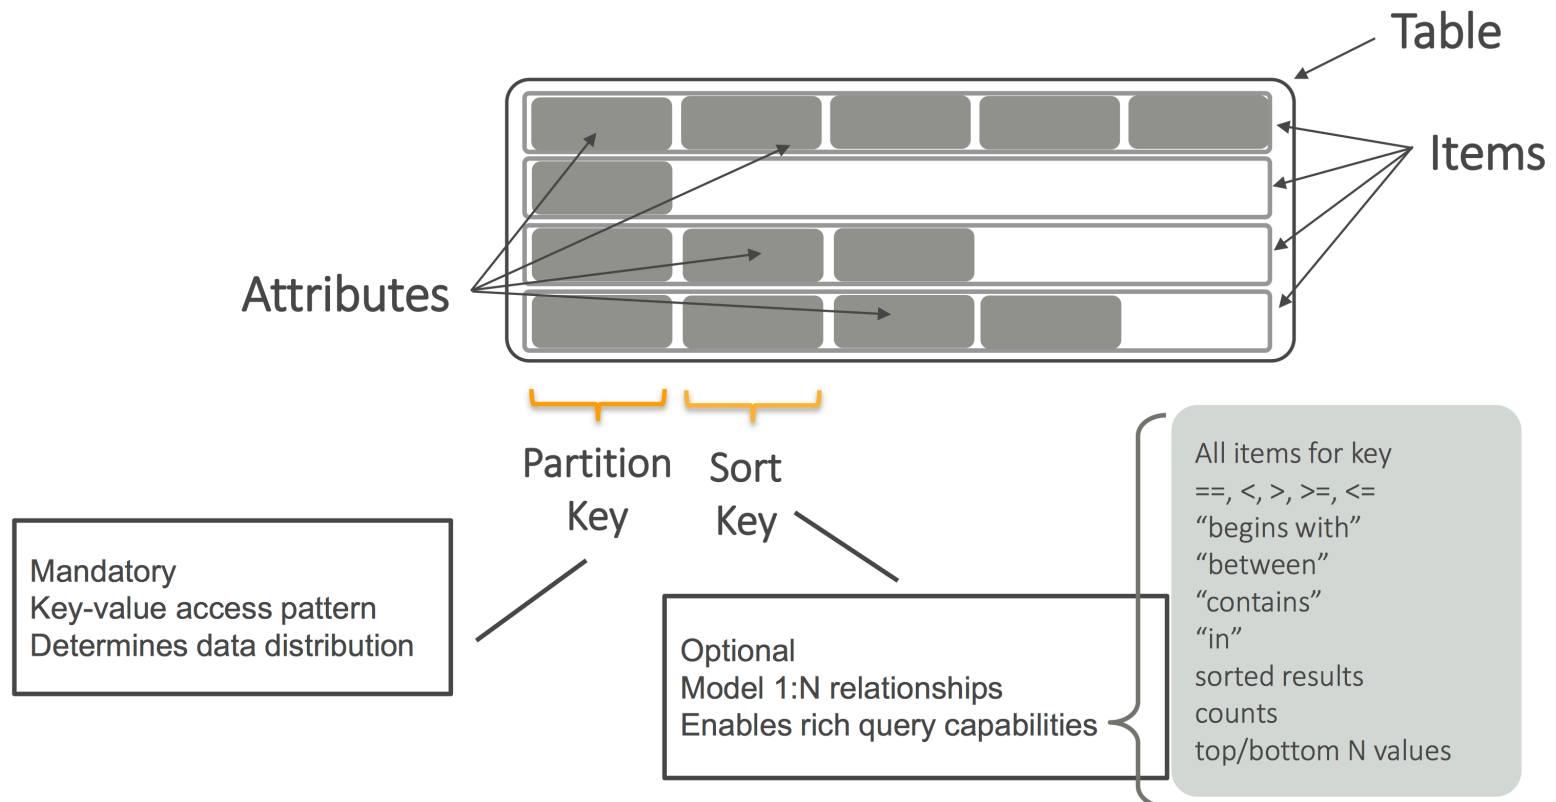
Amazon Dynamo:

- Query model
 - Simple read/write operations on small data items
- ACID properties
 - Weaker consistency model
 - No isolation, mostly single key updates
 - Transactions added recently, but expensive
- Efficiency
 - Tradeoff between performance, cost efficiency, availability and durability guarantees

DynamoDB Data Model

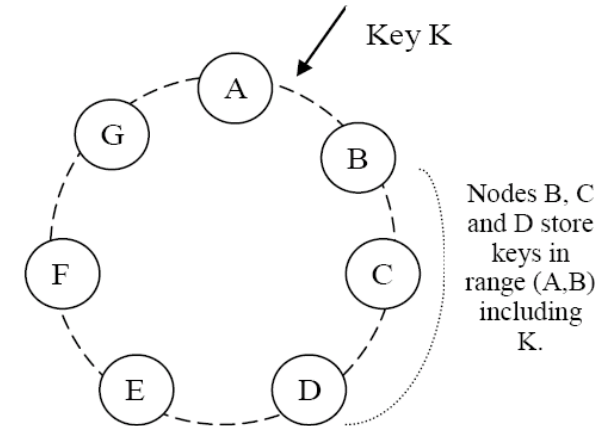
- Tables, items, attributes
 - Table contains multiple items, each item may have multiple attributes
- DynamoDB only requires you specify primary key
 - Unlike relation, attributes and data types for an item are do not have to be defined in advance
 - Total size of item is 400KB
 - Primary key specified when table is created

Core Idea: Table



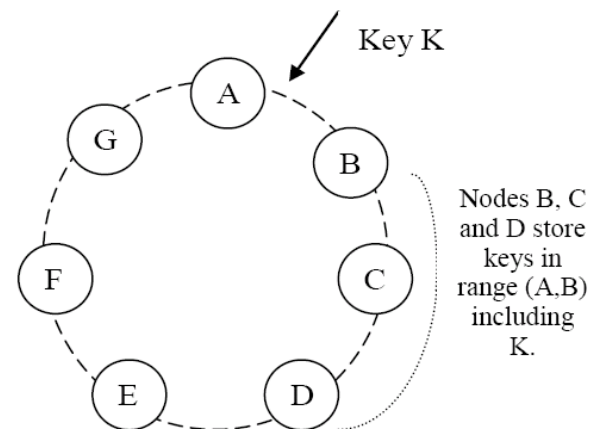
Partition Algorithm

- Consistent hashing: the output range of a hash function is treated as a fixed circular space or “ring”.
- “Virtual Nodes”: Each node can be responsible for more than one virtual node.



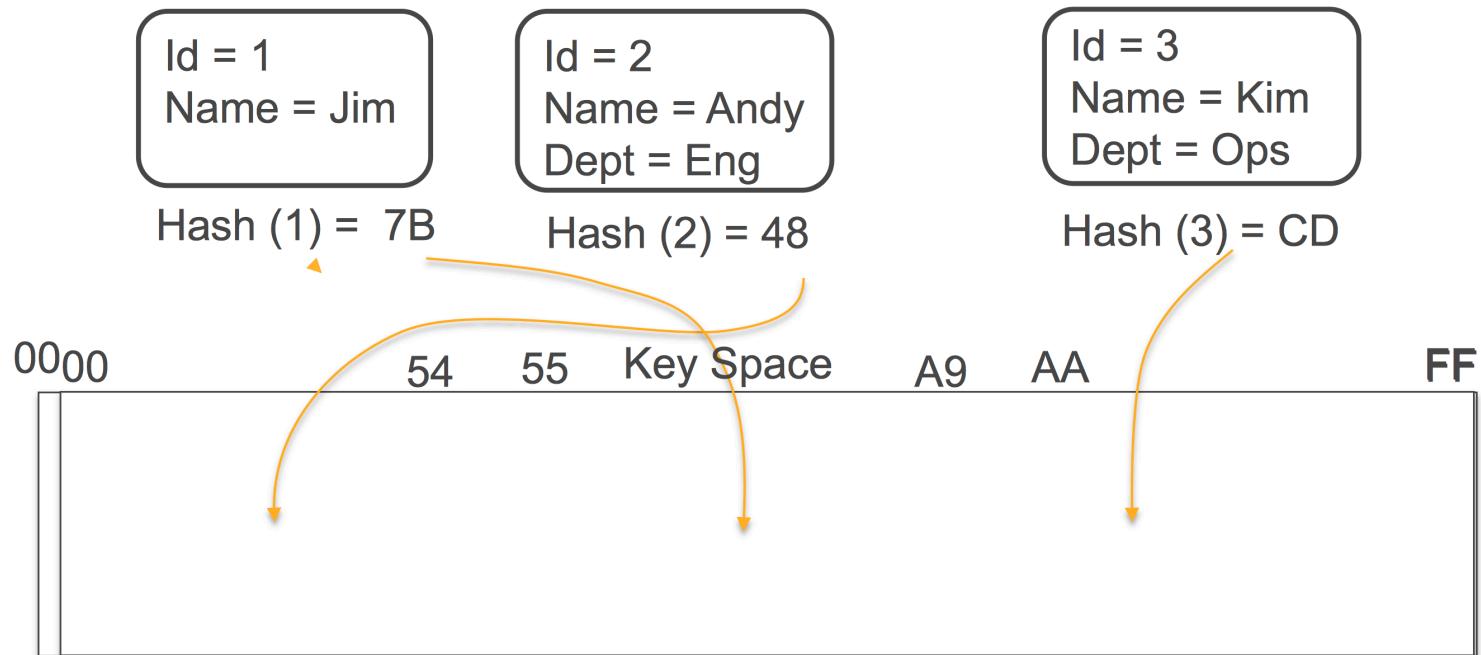
Replication

- Each data item is replicated at N hosts.
- “*preference list*”: The list of nodes that is responsible for storing a particular key.



Partition Table

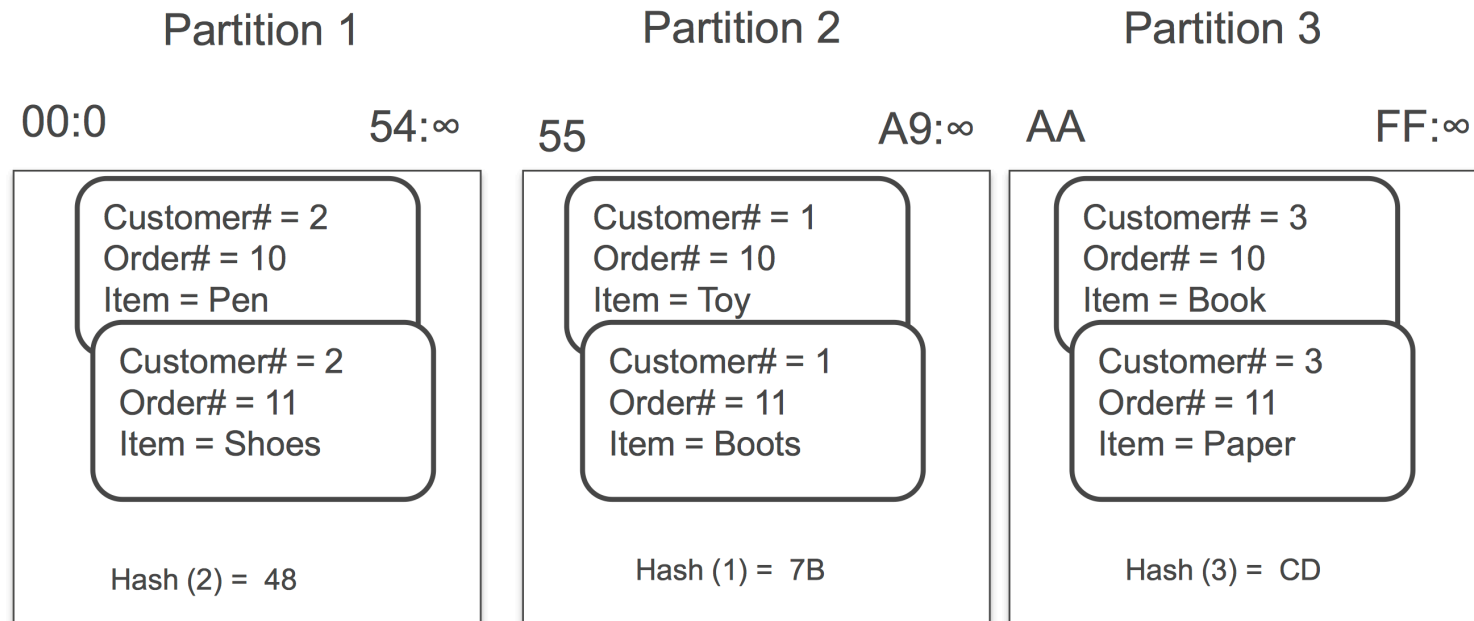
Partition Key uniquely identifies an item
Partition Key is used for building an unordered hash index
Allows table to be partitioned for scale



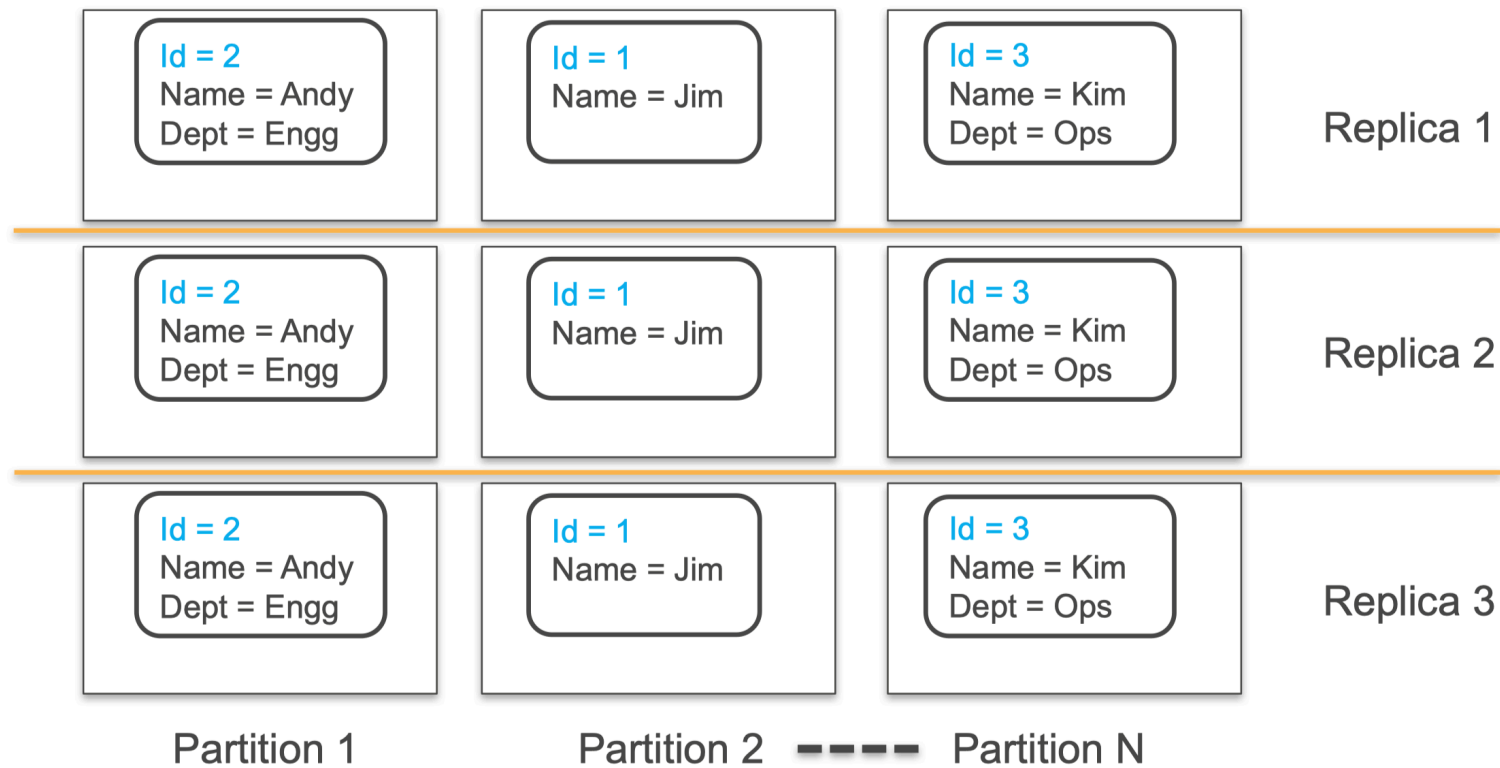
Partition-sort key table

Partition:Sort Key uses two attributes together to uniquely identify an Item
Within unordered hash index, data is arranged by the sort key
No limit on the number of items (∞) per partition key

- Except if you have local secondary indexes



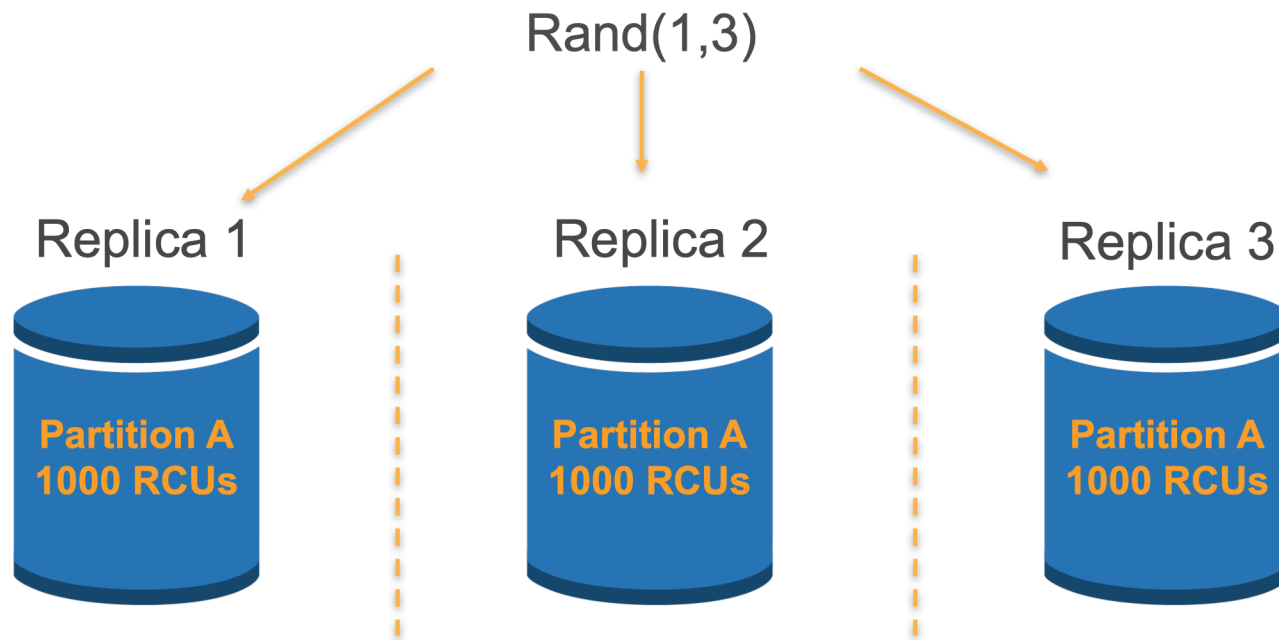
Partitions are Replicated



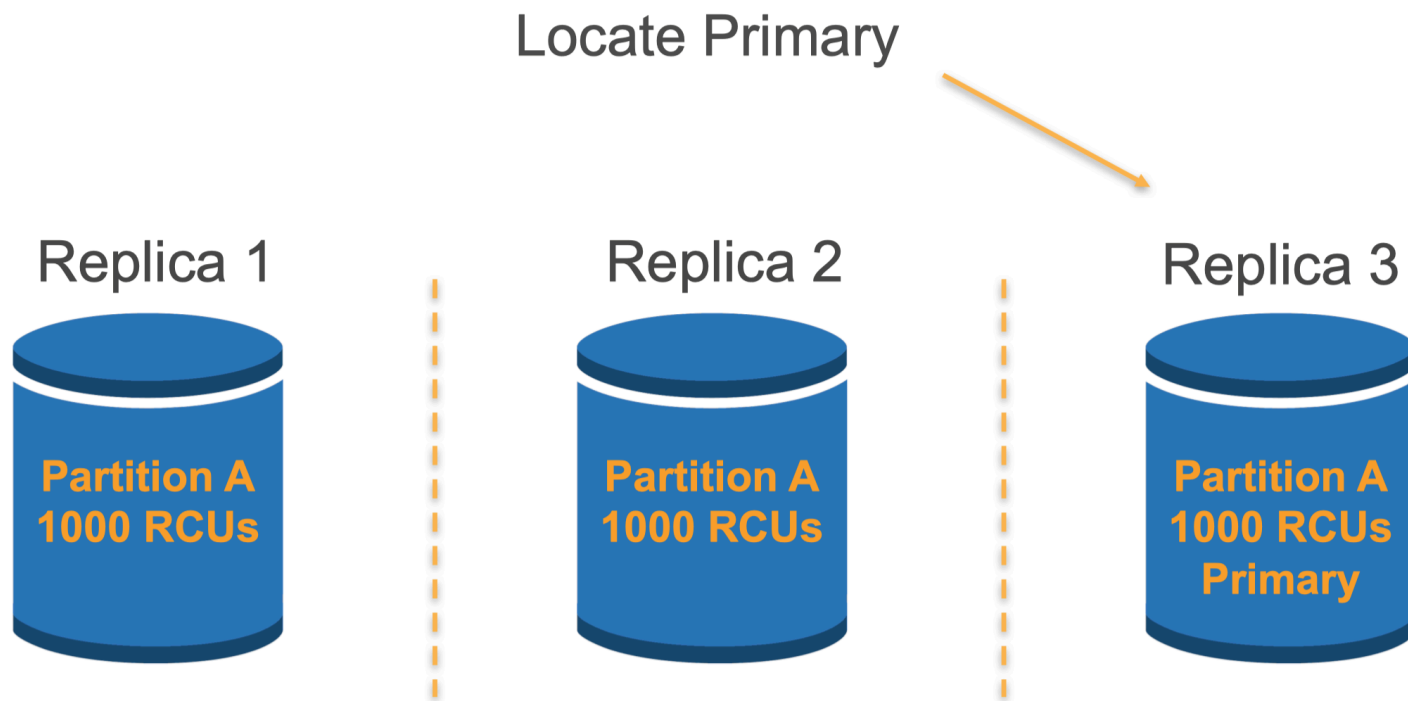
Read Constancy

- Default is eventually consistent
- Can request a strongly consistent value
 - This option may come at cost of reduced availability

Eventually Consistent Reads



Strongly Consistent Reads



Create a table....

- CreateTable API
 - Arguments: {TableName : “Music” ...}
- Every table must have a primary key, which is one of its attributes
- Simple Primary Key
 - Identified by “hash attribute”. DynamoDB uses an unsorted hash index on this value to identify items

```
{  
  TableName: “Music”,  
  KeySchema: { AttributeName: “Artist”, KeyType: “HASH” }  
}
```

Composite Keys

- Use more than one attribute
 - One attribute is used for hash, multiple entries may have the same hash
 - Additional attributes are “ranges” and searched using sorted index, needs to be unique.
- You must provide all of the attributes for a primary key

Create Table Arguments

```
{  
  TableName : "Music",  
  KeySchema: [  
    { AttributeName: "Artist", KeyType: "HASH", },  
    { AttributeName: "SongTitle", KeyType: "RANGE" }  
  ],  
};
```

Can have multiple songs per artist, so first use hash on artist, then use index to find item for song.

```
{ "TableDescription": {
  "TableArn": "arn:aws:dynamodb:us-east1:1289012:table/Music",
  "AttributeDefinitions": [
    { "AttributeName": "Artist", "AttributeType": "S" },
    { "AttributeName": "SongTitle", "AttributeType": "S" } ],
  "ProvisionedThroughput": {...},
  "TableSizeBytes": 0,
  "TableName": "Music",
  "TableStatus": "CREATING",
  "TableId": "12345678-0123-4567-a123-abcdefghijkl",
  "KeySchema": [
    { "KeyType": "HASH", "AttributeName": "Artist" },
    { "KeyType": "RANGE", "AttributeName": "SongTitle" } ],
  "ItemCount": 0, "
  CreationDateTime": 1542397215.37 }
}
```


Table Operations

- CreateTable
- DescribeTable
- DeleteTable
- ListTables

Items

- Items can have multiple attributes
 - Name-value pair
 - Value can be single value, or value set.
- Items are retrieved by a GetItem operation
 - Must provide complete primary key
 - Returns all of the items attributes
 - Eventually consistent

Item Operations

- PutItem
- GetItem
- UpdateItem
- DeleteItem
- BatchGetItem
 - Read up to 100 items from one or more tables.
- BatchWriteItem
 - Create or delete up to 25 items in one or more tables.

GetItem

- Specify key, get back item
 - {"Id":{"N":"1"}}

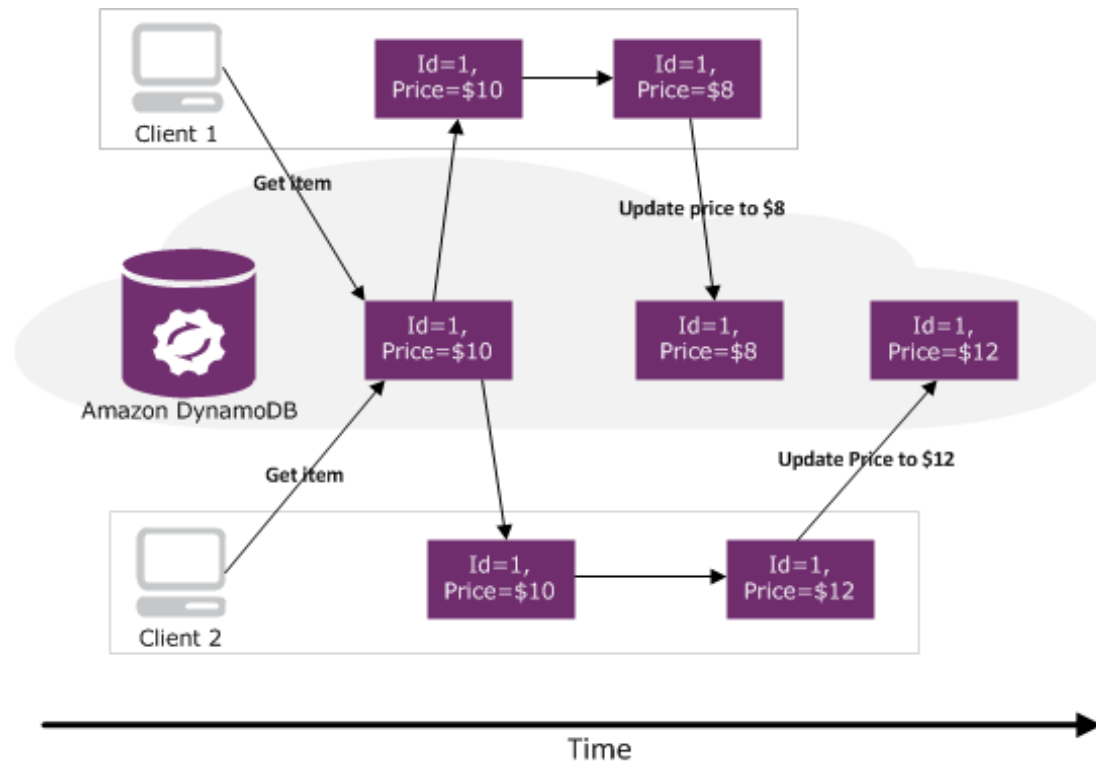
Data Types

- Scalar Types
 - String
 - Unicode with UTF8 encoding
 - Number, binary, boolean, null
- Multi-value Data Types
 - Sets: [“black”, “Green”, “Red”]
 - Documents
 - List: ordered collection of value
 - Map: unordered collection of name-value pairs

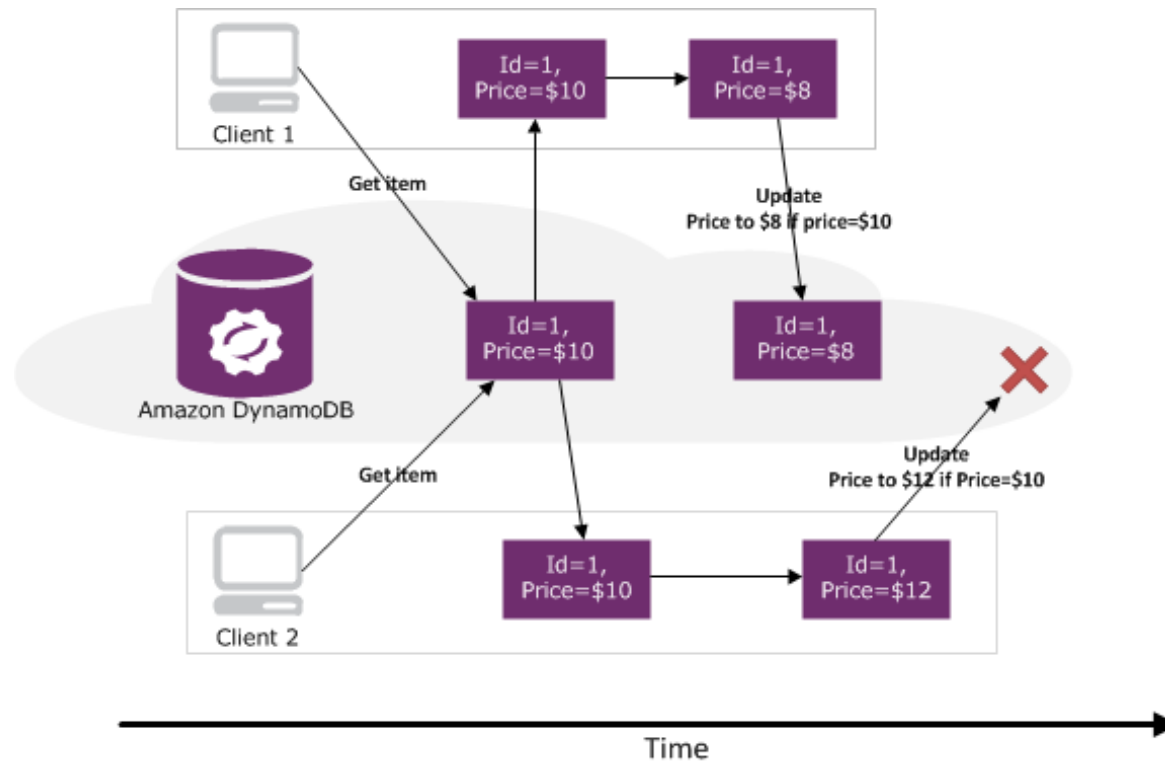
PutItem

```
{  
  "ForumName": {"S": "Amazon DynamoDB"},  
  "Subject": {"S": "New discussion thread"},  
  "Message": {"S": "First post in this thread"},  
  "LastPostedBy": {"S": "fred@example.com"},  
  "LastPostDateTime": {"S": "201603190422"}  
}
```

Concurrent Writes



Conditional Write



Conditional Write

```
{  
  "TableName": "ProductCatalog",  
  Key: {"Id":{"N":"1"}},  
  UpdateExpression "SET Price = :newval" ,  
  "ConditionalExpression": "Price = :currval" ,  
  ExpressionAttributeValues : {  
    ":newval": {"N":"8"},  
    ":currval": {"N":"8"}  
  }  
}
```

Updating an Item

- PutItem
 - Creates a new item, or replaces an existing item
- UpdateItem
 - Changes values of exiting item attributes
- DeleteItem
 - Removes item with matching primary key

Example: Map

```
{  
  Day: "Monday",  
  UnreadEmails: 42,  
  ItemsOnMyDesk: [  
    "Coffee Cup",  
    "Telephone",  
    {  
      Pens: { Quantity : 3},  
      Pencils: { Quantity : 2},  
      Erasers: { Quantity : 1}  
    }  
  ]  
}
```

Specifying Attributes

- Top level Attributes
 - Day, UnreadEmails
- Nested Attributes – [n] for list, “.” for map elements.
 - ItemsOnMyDesk[1]
 - .ItemsOnMyDesk[3].Pens.Quantity

Projection Expressions

- You can project a subset of attributed by listing them:
{
 "TableName": "ProductCatalog",
 "Key": {"Id": {"N": "123"}},
 "ProjectionExpression":
 "Description, RelatedItems[0], ProductReviews.FiveStar"
}

Attribute Names and Values

- Cannot include many names or values directly, so use a dictionary.
- Names: GET {“TableName”: “ProductCatalog”,
“Key “: {“Id”:{“N”:“123”}} ,
“ProjectionExpression”: “#pr.#1star” ,
“ExpressionAttributeNames”:
 {“#pr”:“ProductReviews”, “#1star”:“OneStar”}
}
- Values: { “:c”: { “S”: “Black” }, “:p”: { “N”: “500” } }

Query Operation

- Query uses primary key to find items in a table
 - Only works on one table, no joins!!
- By default all attributes are returned, but you can do a projection by specifying a *ProjectionExpression* parameter
- Can return a subset of items by using a *KeyConditionExpression* on a range attribute

What about Queries?

- To build a DynamoDB Database, we need to know the questions
- Eg:
 - Find hotels in a given area
 - Find information about a given hotel
 - Find points of interest near a given hotel

Consider this example data

Table Name	Primay Key Type	Hash Attribute Name and Type	Range Attribute Name and Type
ProductCatalog(<u>Id</u> , ...)	Hash	Attribute Name: Id Type: Number	-
Forum(<u>Name</u> , ...)	Hash	Attribute Name: Name Type: String	-
Thread(<u>ForumName</u> , <u>Subject</u> , ...)	Hash and Range	Attribue Name: ForumName Type: String	Attribue Name: Subject Type: String
Reply(<u>Id</u> , <u>ReplyDateType</u> , ...)	Hash and Range	Attribue Name: Id Type: String	Attribute Name; ReplyDateType Type: String

Queries

- Query the Thread table for a particular ForumName (hash attribute). All of the items with that ForumName value will be read by the query, because the range attribute (Subject) is not included in *KeyConditionExpression*.
 - ForumName = :name
- Query the Thread table for a particular ForumName (hash attribute), but this time return only the items with a given Subject (range attribute).
 - Forum = :name and Subject = :subj
- Query the Reply table for a particular Id (hash attribute), but return only those items whose ReplyDateTime (range attribute) begins with certain characters.
 - Id = :id and begins_with(ReplyDateTime, :dt)

Scan

- Read every item in the table
- Return every data item
 - Can use *ProjectionExpression* parameter to return only subset of columns.

What about “where” clause

- Can filter results of Query or Scan using *FilterExpression* parameter.

Filter...

- Query the Thread table for a particular ForumName (hash attribute) and Subject (range attribute). Of the items that are found, return only the most popular discussion threads—for example, those threads with more than a certain number of Views (views is a reserved word in DynamoDB, so we need to “escape” it).
 - #V > :num
- Scan the Thread table and return only the items that were last posted to by a particular user.
 - LastPostedBy = :name

Secondary Indexes

- Items in table consist of a hash key, range key and attributes
- Fast lookup on hash key or hash and range
 - Not on attribute values or range alone
- What if you want to find a value quickly?
 - Global secondary index comes to the rescue
 - Specify alternative key/range attributes

Global secondary index (GSI)

Alternate partition and/or sort key

Index is across all partition keys

Online indexing

Table



RCUs/WCUs
provisioned separately
for GSIs

GSIs



KEYS_ONLY



INCLUDE A3

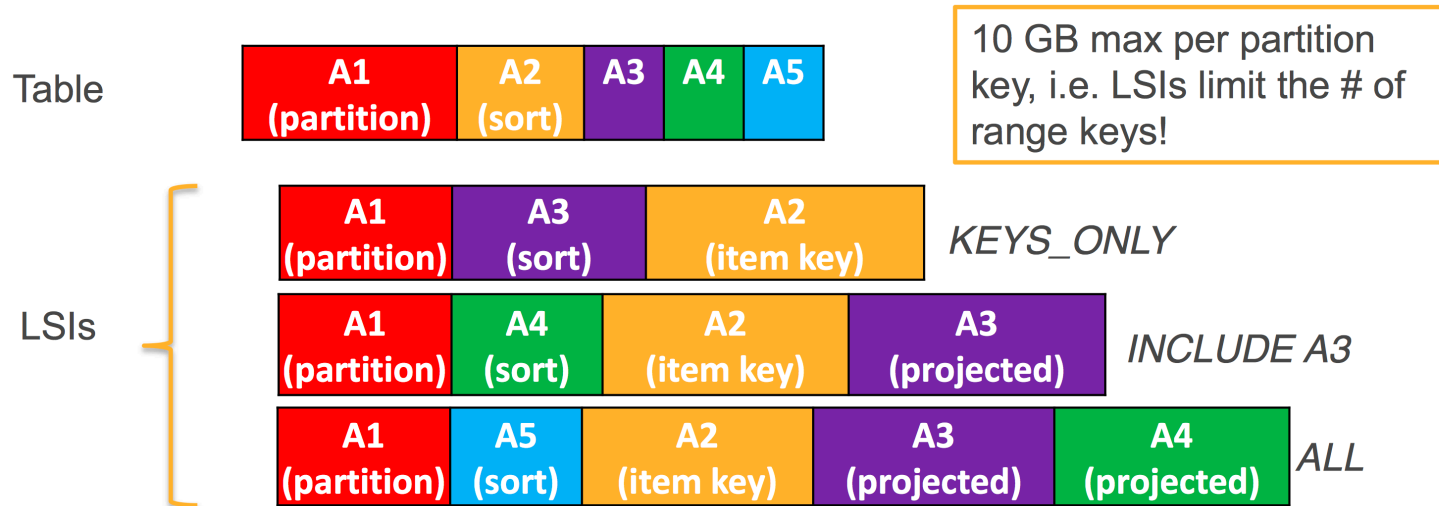


ALL

Local secondary index (LSI)

Alternate sort key attribute

Index is local to a partition key



GameScores

UserId (Hash key)	GameTitle (Range key)	TopScore	TopScoreDateTime	Wins	Losses	
"101"	"Galaxy Invaders"	5842	"2013-09-15:17:24:31"	21	72	...
"101"	"Meteor Blasters"	1000	"2013-10-22:23:18:01"	12	3	...
"101"	"Starship X"	24	"2013-08-31:13:14:21"	4	9	...
"102"	"Alien Adventure"	192	"2013-07-12:11:07:56"	32	192	...
"102"	"Galaxy Invaders"	0	"2013-09-18:07:33:42"	0	5	...
"103"	"Attack Ships"	3	"2013-10-19:01:13:24"	1	8	...
"103"	"Galaxy Invaders"	2317	"2013-09-11:06:53:00"	40	3	...
"103"	"Meteor Blasters"	723	"2013-10-19:01:13:24"	22	12	...
"103"	"Starship X"	42	"2013-07-11:06:53:00"	4	19	...
...	

From W. Vogels)

What about

- Query by user is fast, but
 - What is the top score ever recorded for the game "Meteor Blasters"?
 - Which user had the highest score for "Galaxy Invaders"?
 - What was the highest ratio of wins vs. losses?
- Would have to use a scan to do this

Secondary Global Index

GameTitleIndex

<i>GameTitle</i> (Hash key)	<i>TopScore</i> (Range key)	<i>UserId</i>
"Alien Adventure"	192	"102"
"Attack Ships"	3	"103"
"Galaxy Invaders"	0	"102"
"Galaxy Invaders"	2317	"103"
"Galaxy Invaders"	5842	"101"
"Meteor Blasters"	723	"103"
"Meteor Blasters"	1000	"101"
"Starship X"	24	"101"
"Starship X"	42	"103"
...

Many to Many

Table	Primary Key		Data Attributes...		
	Partition Key	Sort Key (and GSI PK)			
	Invoice-92551	Inv_ID: Invoice-92551 (invoice ID)	Dated: 2018-02-07 (date created)	More attributes of this invoice...	
		Bill_ID: Bill-4224663 (bill ID)	Dated: 2017-12-03 (date created)	Attributes of this bill in this invoice..	
		Bill_ID: Bill-4224687 (bill ID)	Dated: 2018-01-09 (date created)	Attributes of this bill in this invoice..	
	Invoice-92552	Inv_ID: Invoice-92552 (invoice ID)	Dated: 2018-03-04 (date created)	More attributes of this invoice...	
		Bill_ID: Bill-4224687 (bill ID)	Dated: 2018-01-09 (date created)	Attributes of this bill in this invoice..	
	Bill-4224663	Bill_ID: Bill-4224663 (bill ID)	Dated: 2017-12-03 (date created)	More attributes of this bill...	
	Bill-4224687	Bill_ID: Bill-4224687 (bill ID)	Dated: 2018-01-09 (date created)	More attributes of this bill...	

GSI	Primary Key	Projected Attributes...	
	Partition Key		
	Bill-4224663	Bill_ID: Bill-4224663 (table primary key)	Attributes of this bill...
		Inv_ID: Invoice-92551 (table primary key)	Attributes of this bill <i>in this invoice</i> ..
	Bill-4224687	Bill_ID: Bill-4224687 (table primary key)	Attributes of this bill...
		Inv_ID: Invoice-92551 (table primary key)	Attributes of this bill <i>in this invoice</i> ..
		Inv_ID: Invoice-92552 (table primary key)	Attributes of this bill <i>in this invoice</i> ..
	Invoice-92551	Inv_ID: Invoice-92551 (table primary key)	Attributes of this invoice...
	Invoice-92552	Inv_ID: Invoice-92552 (table primary key)	Attributes of this invoice...